

Ray Tracing for Games

Dr. Jacco Bikker - IGAD/BUAS, Breda, February 7

Welcome!

11





Thursday
09:00 – 14:00

advanced Whitted
audio, AI & physics
faster Whitted
Heaven7

Wednesday
13:00 – 17:00

course intro
LH2
template
Whitted
refactoring
RT-centric games

LAB 1

LAB 2

work @ home

End result day 2:

A solid Whitted-style
ray tracer, as a basis
for subsequent work.

Friday
09:00 – 17:00

optimization
profiling, rules of
engagement
threading



LAB 3

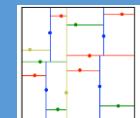
SIMD
applied SIMD
SIMD triangle
SIMD AABB

LAB 4

End result day 3:
A 5x faster tracer.

Monday
09:00 – 17:00

acceleration
grid, BVH, kD-tree
SAH
binning



LAB 5

refitting
top-level BVH
threaded building

LAB 6

End result day 4:
A real-time tracer.

Tuesday
09:00 – 17:00

Monte-Carlo
Cook-style
glossy, AA
area lights, DOF



LAB 7

path tracing

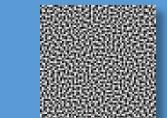


LAB 8

End result day 5:
Cook or Kajiya.

Thursday
09:00 – 17:00

random numbers
stratification
blue noise



LAB 9

importance
sampling
next event
estimation

LAB 10

End result day 6:
Efficiency.

Friday
09:00 – 17:00

future work



LAB 11

path guiding



LAB 10

End result day 6:
Great product.

GLOBAL GAME JAM

Agenda:

- Filtering
- Grand Recap
- Future Work



Reducing Noise in Path Tracing

We can increase ‘exposure’ in a pathtracer by taking more samples.

Or we can use Importance Sampling.

```
rics
  & (depth < MAXDEPTH)
  c = inside ? 1 : 1.2;
  nt = nc / ncy ddn = ddn * c;
  os2t = 1.0f - nnt * nnt;
  D, N );
  )

at a = nt - nc, b = nt + nc;
at Tr = 1 - (R0 + (1 - R0) *
Tr) R = (D * nnt - N * (ddn
E * diffuse;
= true;

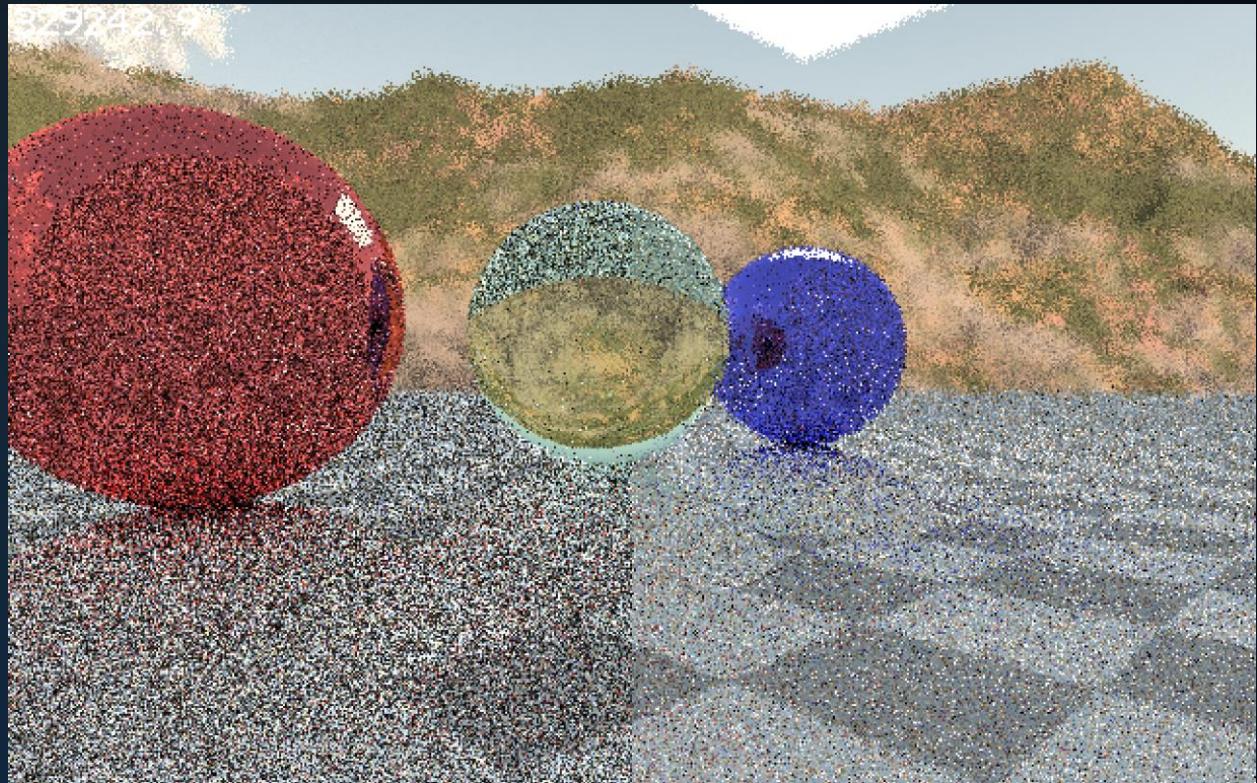
  refl + refr)) && (depth < MAXDEPTH)

D, N );
refl * E * diffuse;
= true;

MAXDEPTH)

survive = SurvivalProbability( diffuse
estimation - doing it properly, closely
df;
radiance = SampleLight( &rand, I, &L, &lightbuf,
e.x + radiance.y + radiance.z) > 0) && (dot( N
e = true;
at brdfPpdf = EvaluateDiffuse( L, N ) * Psurvive;
at3 factor = diffuse * INVPI;
at weight = Mis2( directPpdf, brdfPpdf );
at cosThetaOut = dot( N, L );
E * (weight * cosThetaOut) / directPpdf) * (radiance
random walk - done properly, closely following Smits
alive)

at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf
survive;
pdf;
in = E * brdf * (dot( N, R ) / pdf);
ision = true;
```



```
rics
  & (depth < MAXDEPTH)
  c = inside ? 1 : 1.2;
  nt = nc / ncy ddn = d * d;
  os2t = 1.0f - nnt * ddn;
  D, N );
  )
}

at a = nt - nc, b = nt + nc;
at Tr = 1 - (R0 + (1 - R0) * Tr) R = (D * nnt - N * (ddn
E * diffuse;
= true;

  refl + refr)) && (depth < MAXDEPTH)

  D, N );
  refl * E * diffuse;
  = true;

MAXDEPTH)

survive = SurvivalProbability( diffuse );
estimation = doing it properly, closely
df;
radiance = SampleLight( &rand, I, &L, &lightbuf,
e.x + radiance.y + radiance.z) > 0) && (dot( N

  v = true;
  at brdfPpdf = EvaluateDiffuse( L, N ) * Psurvive;
  at3 factor = diffuse * INVPI;
  at weight = Mis2( directPpdf, brdfPpdf );
  at cosThetaOut = dot( N, L );
  E * (weight * cosThetaOut) / directPpdf) * (radiance
random walk - done properly, closely following Smith's
alive);

at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf );
survive;
pdf;
  E * brdf * (dot( N, R ) / pdf);
  = true;
```

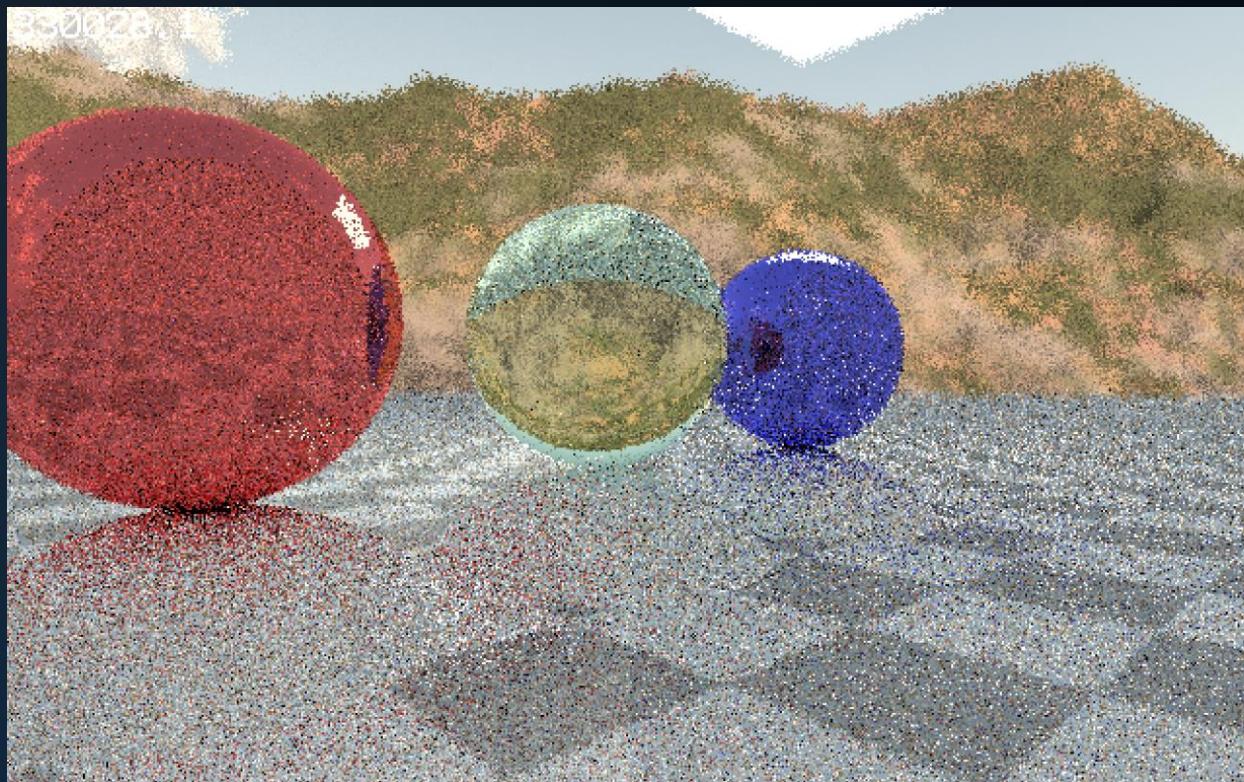
Reducing Noise in Path Tracing

We can increase ‘exposure’ in a pathtracer by taking more samples.

Or we can use Importance Sampling.

Or Next Event Estimation.

However, at interactive rates it will be quite hard to reach a converged image.



Reducing the Problem - Filtering

Core idea:

Exploit the fact that illumination is typically low-frequent:

Nearby pixels tend to converge to similar values, so we should be able to use information gathered for one pixel to improve the estimate of the next.

Note:

Unless neighboring pixels actually converge to the same value, filtering will introduce bias.

```
rics
  & (depth < MAXDEPTH)
  c = inside ? 1 : 1.2f;
  nt = nc / c;
  os2t = 1.0f - nnt * nnt;
  D, N );
  R = (D * nnt - N * (dd
E * diffuse;
= true;

-
refl + refr)) && (depth < MAXDEPTH)

  D, N );
  refl * E * diffuse;
= true;

MAXDEPTH)

survive = SurvivalProbability( diffuse )
estimation - doing it properly, closely
if;
radiance = SampleLight( &rand, I, &L, &lightbuf,
e.x + radiance.y + radiance.z) > 0) && (dot( N
e = true;
at brdfPpdf = EvaluateDiffuse( L, N ) * Psurvive;
at3 factor = diffuse * INVPI;
at weight = Mis2( directPpdf, brdfPpdf );
at cosThetaOut = dot( N, L );
E * ((weight * cosThetaOut) / directPpdf) * (radiance
random walk - done properly, closely following Smillie
ive)

;
at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf
survive;
pdf;
n = E * brdf * (dot( N, R ) / pdf);
ision = true;
```



Filter kernels

For the actual filtering, we apply a kernel.

```

Pixel FilteredValue( i_x, i_y, halfWidth )
    sum = 0
    summedWeight = 0
    for j_x = i_x - halfWidth to i_x + halfWidth
        for j_y = i_y - halfWidth to i_y + halfWidth
            sum += ReadPixel( j_x, j_y ) * weight( j_x, j_y )
            summedWeight += weight( j_x, j_y )
    return sum / summedWeight

```

$$\hat{c}_i = \frac{\sum_{j \in \mathcal{N}_i} c_j w(i, j)}{\sum_{j \in \mathcal{N}_i} w(i, j)}$$



Filter kernels

For the actual filtering, we apply a kernel.

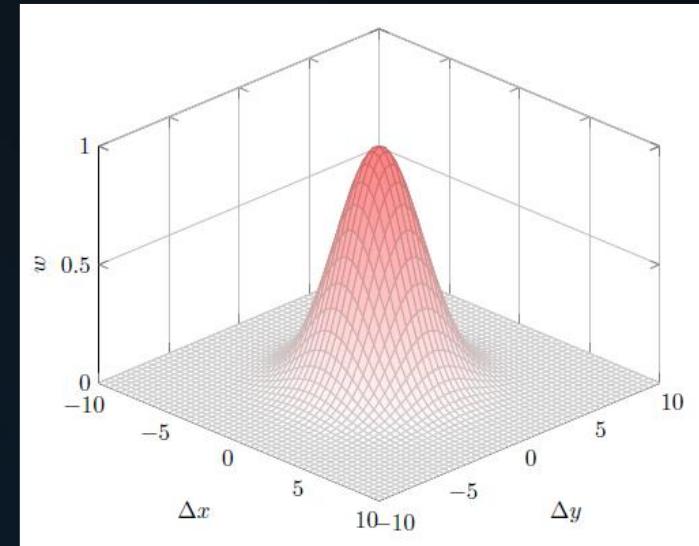
```

Pixel FilteredValue( i_x, i_y, halfWidth )
    sum = 0
    summedWeight = 0
    for j_x = i_x - halfWidth to i_x + halfWidth
        for j_y = i_y - halfWidth to i_y + halfWidth
            sum += ReadPixel( j_x, j_y ) * weight( j_x, j_y )
            summedWeight += weight( j_x, j_y )
    return sum / summedWeight

```

Here, **weight** or w is the weight function. We could simply use the Gaussian kernel:

$$w(i, j) = \exp\left(\frac{-\|p_i - p_j\|^2}{2\sigma_d^2}\right),$$

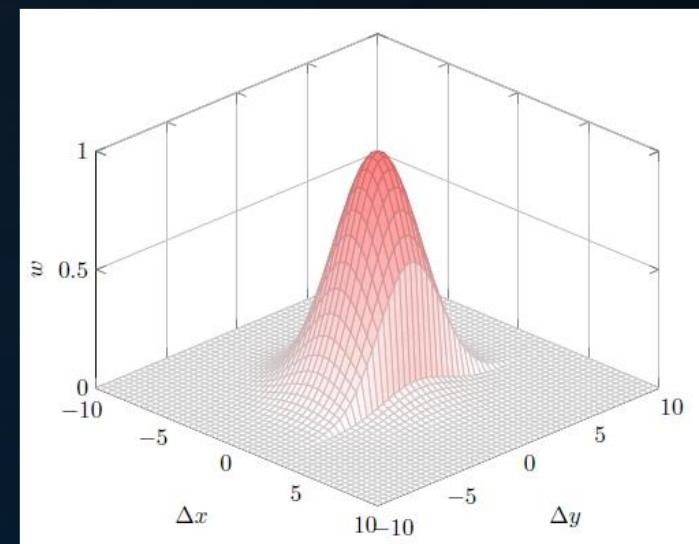
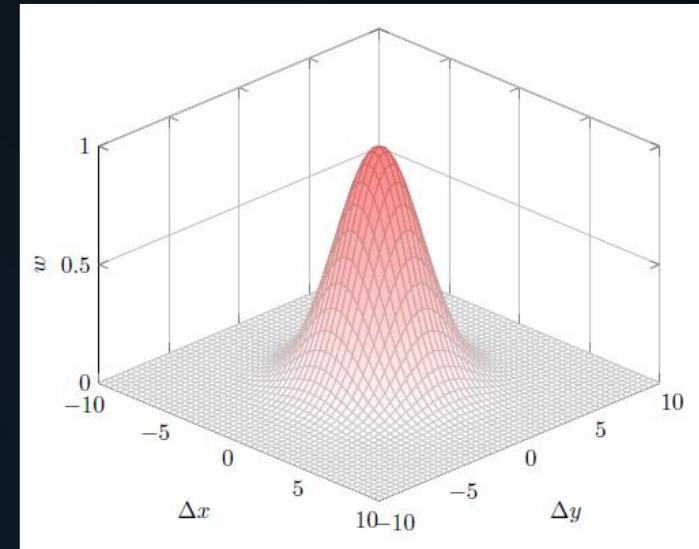
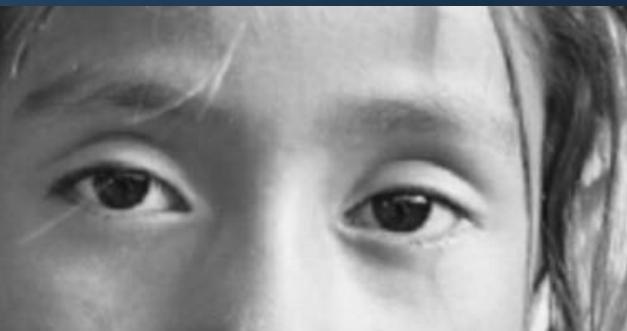


Filter kernels

A Gaussian filter (as well as other low-pass filters) blurs out high frequency details.

We can improve on this using a non-linear bilateral filter*.

$$w(i, j) = \exp\left(\frac{-\|p_i - p_j\|^2}{2\sigma_d^2}\right) \times \exp\left(\frac{-\|c_i - c_j\|^2}{2\sigma_r^2}\right)$$



*: Tomasi & Manduchi, Bilateral filtering for gray and color images. ICCV '98.



Adding this separation to an existing renderer:

- store albedo at the primary intersection (simple material property);
- at the end of the pipeline: illumination = sample / max(epsilon, albedo).



```
rics
  & (depth < MAXDEPTH)
  c = inside ? 1 : 1.2;
  nt = nc / ncy ddn = ddn * c;
  os2t = 1.0f - nnt * os2t;
  D, N );
}
)

at a = nt - nc, b = nt + nc;
at Tr = 1 - (R0 + (1 - R0) *
Tr) R = (D * nnt - N * fdd
E * diffuse;
= true;

-
at refl + refr) && (depth < MAXDEPTH)
D, N );
at refl * E * diffuse;
= true;

MAXDEPTH)

survive = SurvivalProbability( diffuse );
estimation - doing it properly, closely
if;
radiance = SampleLight( &rand, I, &L, &lightbuf,
e.x + radiance.y + radiance.z ) > 0) && (dot( N
e = true;
at brdfPpdf = EvaluateDiffuse( L, N ) * Psurvive;
at t3 factor = diffuse * INMPI;
at weight = Mis2( directPpdf, brdfPpdf );
at cosThetaOut = dot( N, L );
E * (weight * cosThetaOut) / directPpdf) * (radiance
random walk - done properly, closely following Saini
alive);

at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf
survive;
pdf;
E * brdf * (dot( N, R ) / pdf);
ision = true;
```



Example of denoising using a cross-bilateral filter in Blender.



sinmantyx.wordpress.com/2015/12/02/denoising-cycles-renders-in-blender-theory



Reprojection

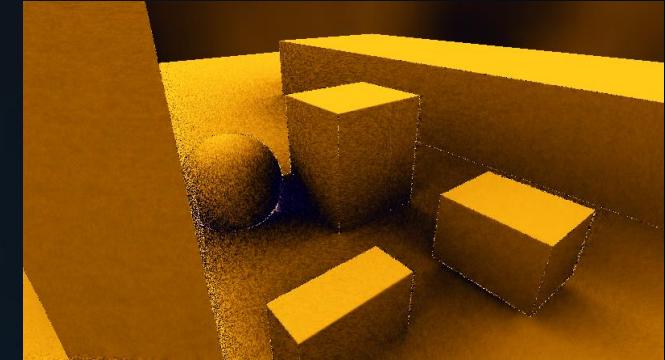
Core idea:

In an animation, samples taken for the previous frame are meaningful for the current frame. *We can supply the filter with more data by looking back in time.*



Reprojection

Core idea:



In an animation, samples taken for the previous frame are meaningful for the current frame. *We can supply the filter with more data by looking back in time.*

Problem: in an animation, the camera and/or the geometry moves. We need to find the location of a pixel in the previous frame(s).

Solution: use the camera matrices.

$$M_{4 \times 4} \begin{pmatrix} x_{world} \\ y_{world} \\ z_{world} \\ 1 \end{pmatrix} = \begin{pmatrix} x_{screen} \\ y_{screen} \\ z_{screen} \\ 1 \end{pmatrix} \rightarrow M_{4 \times 4}^{-1} \begin{pmatrix} x_{screen} \\ y_{screen} \\ z_{screen} \\ 1 \end{pmatrix} = \begin{pmatrix} x_{world} \\ y_{world} \\ z_{world} \\ 1 \end{pmatrix}$$

(finally, apply the matrix of the previous frame to obtain the screen location in the previous frame.)



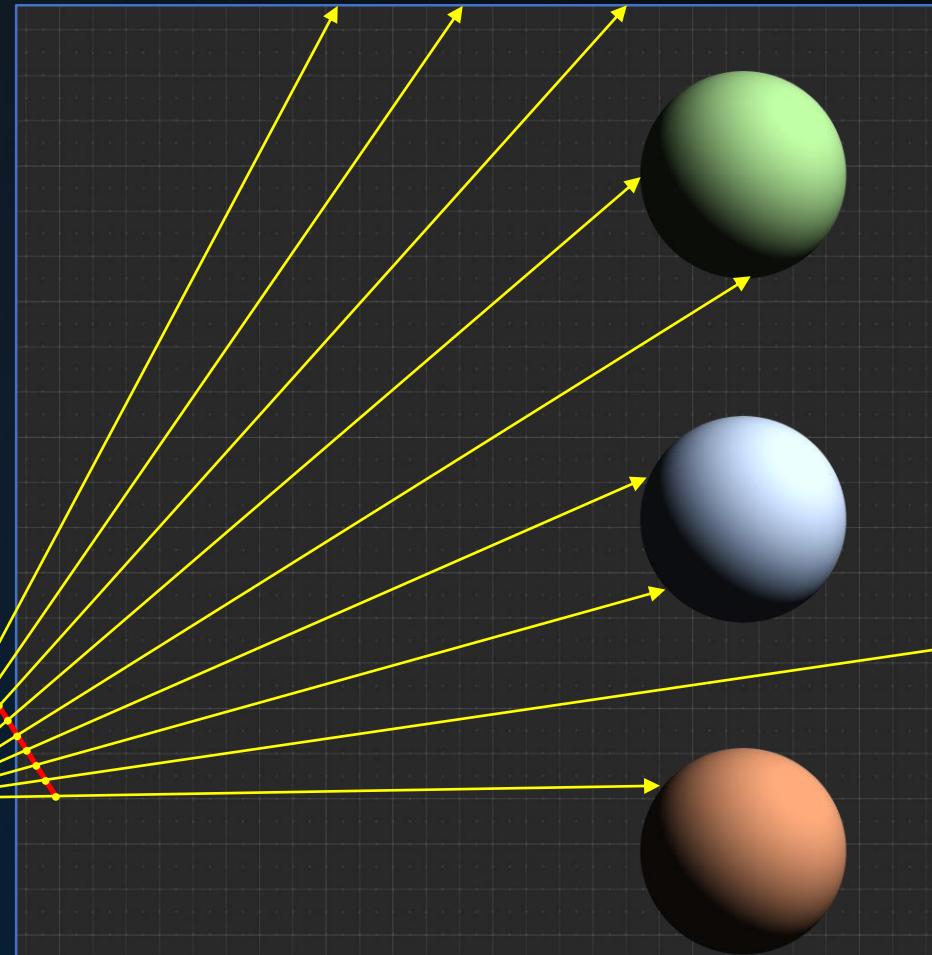
Agenda:

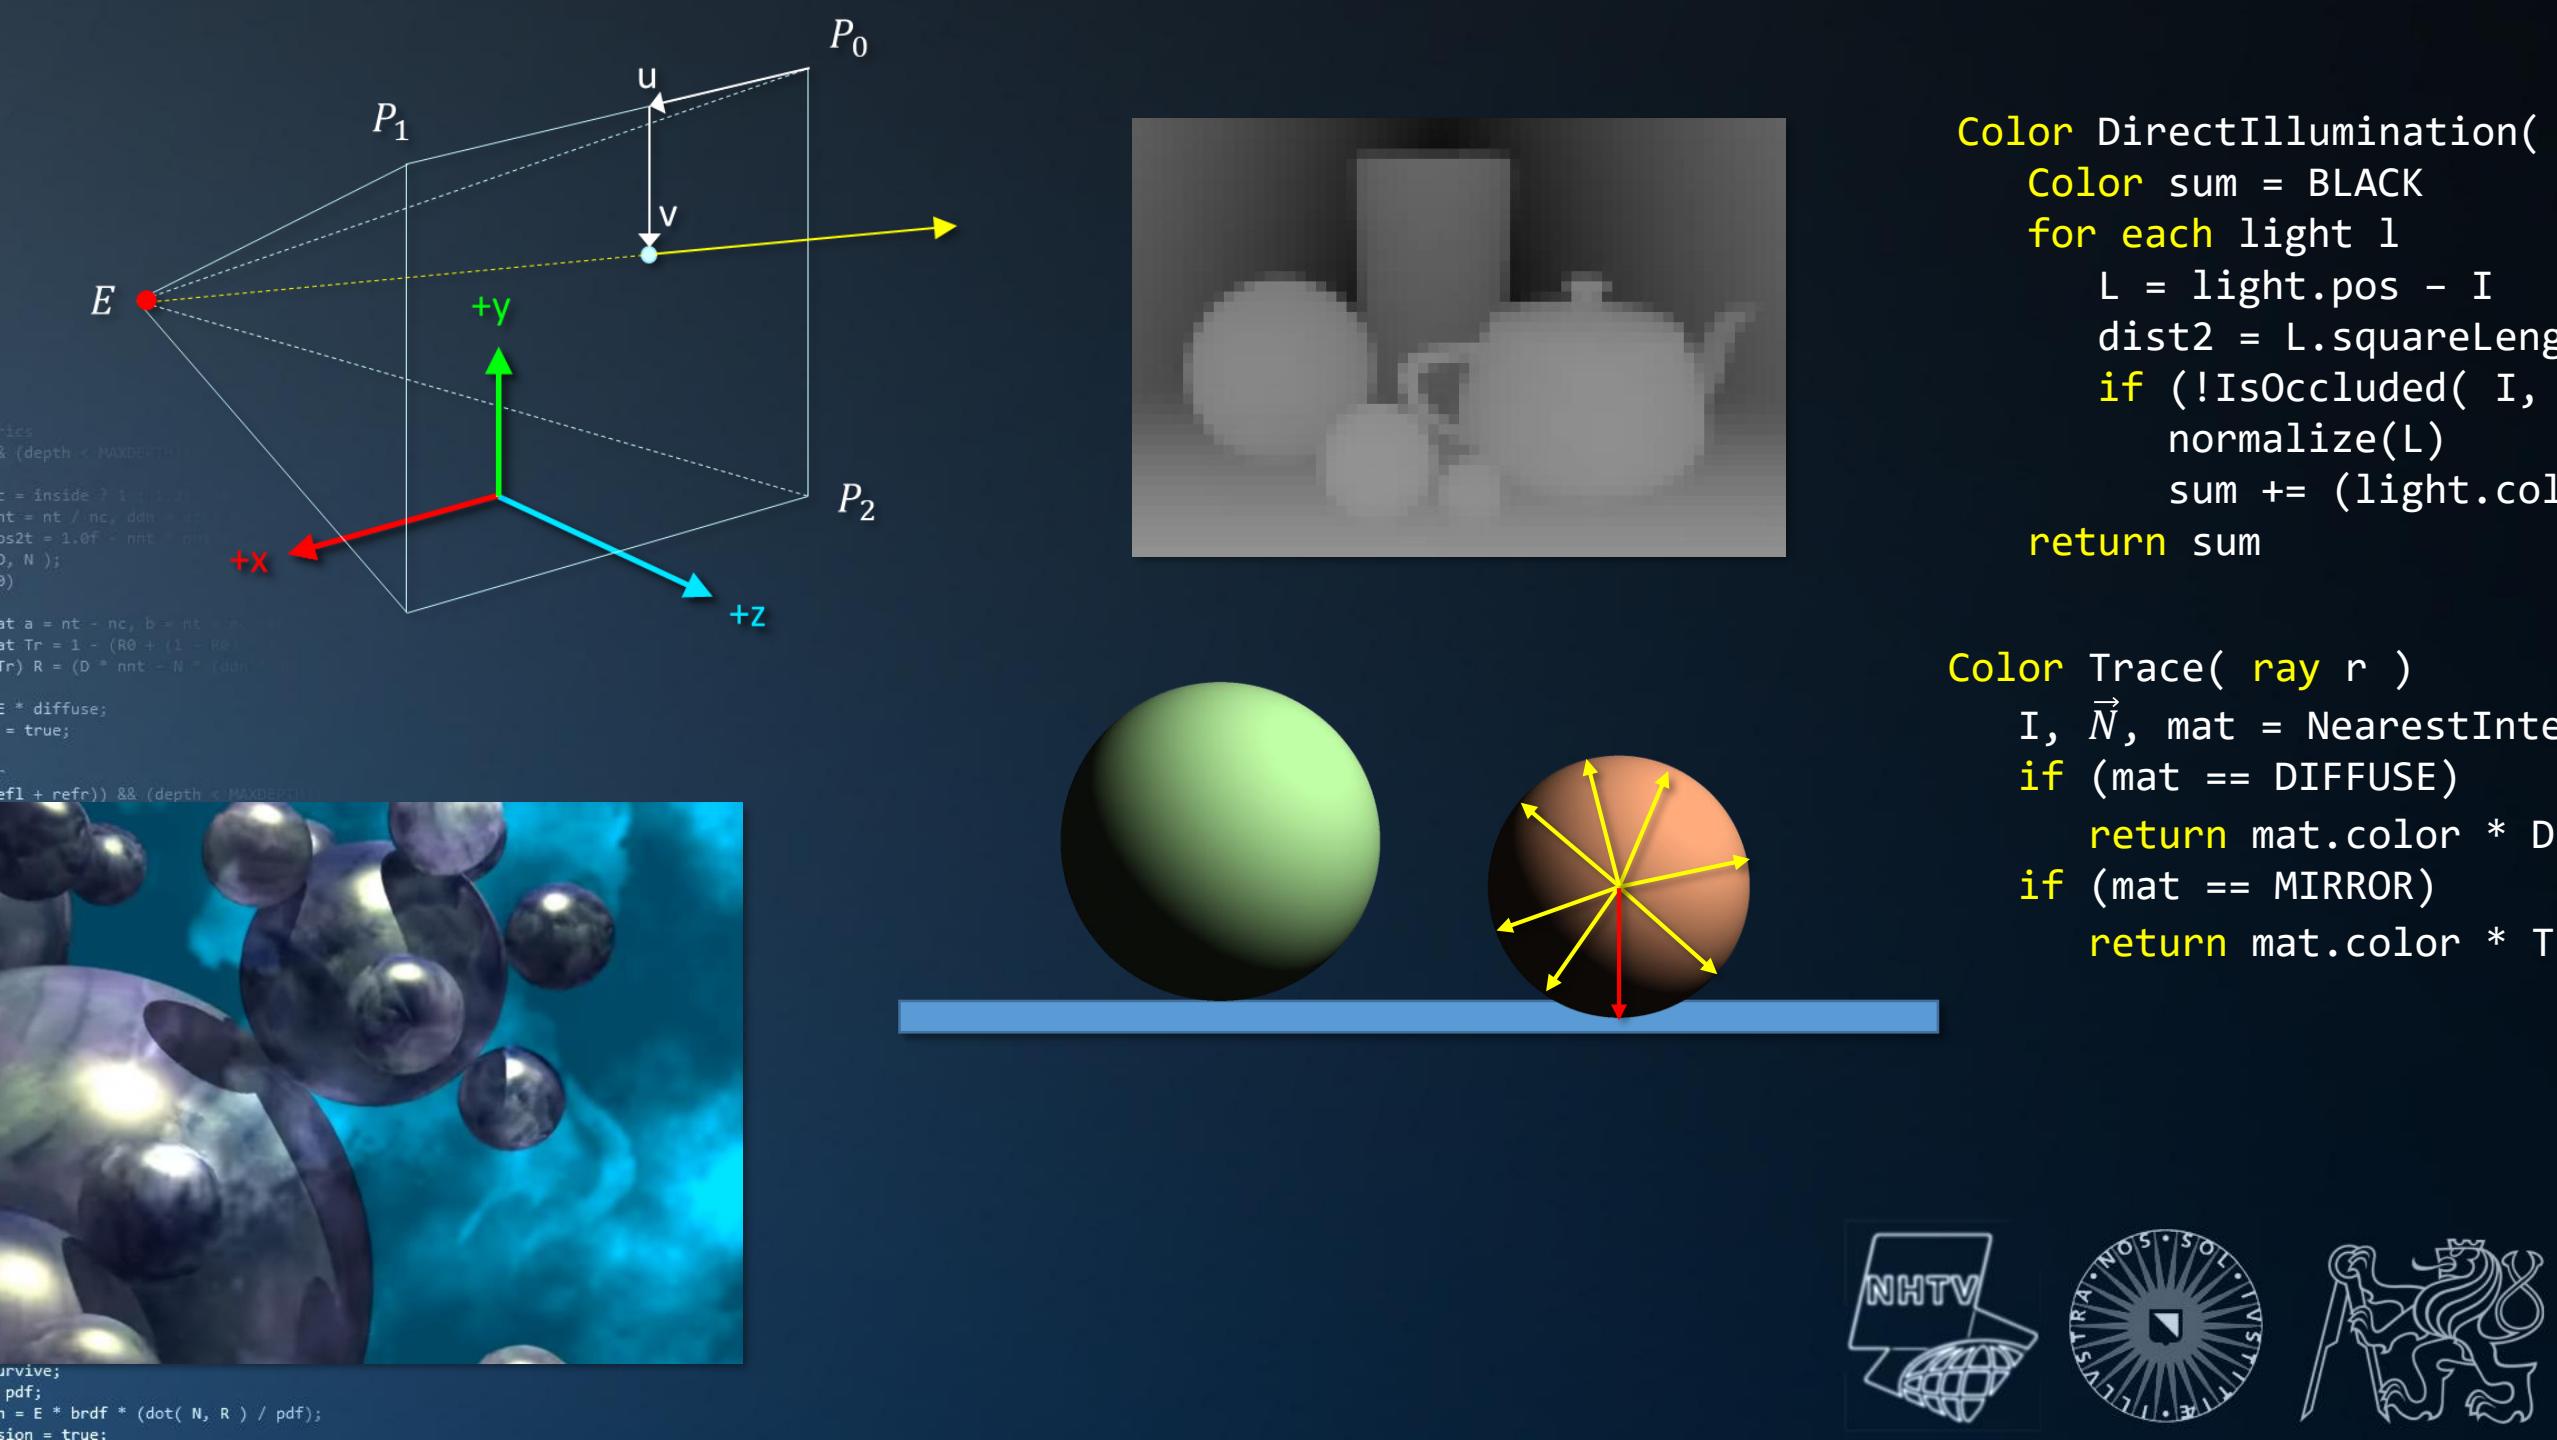
- Filtering
- Grand Recap
- Future Work

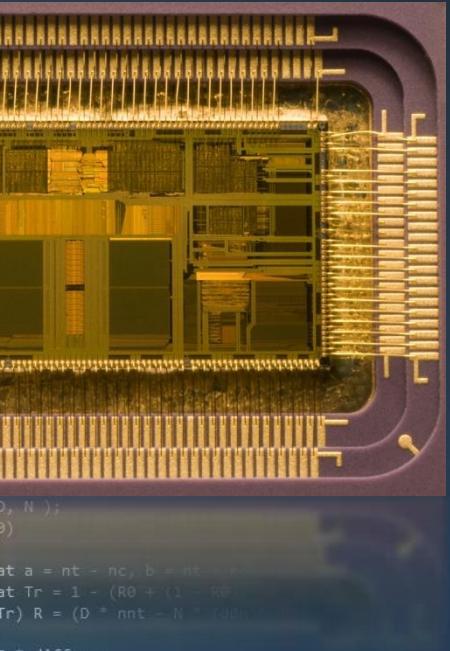




```
E * ((weight * cosThetaOut) / directPdf) * (radiance  
random walk - done properly, closely following Smits  
ive)  
;  
at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf  
survive;  
pdf;  
in = E * brdf * (dot( N, R ) / pdf);  
ision = true;
```







```

    ...
    at a = nt - nc, b = nt - nc;
    at Tr = 1 - (R0 + (1 - R0) * Tr);
    Tr = (D * nnt - N * Tdd);
    ...

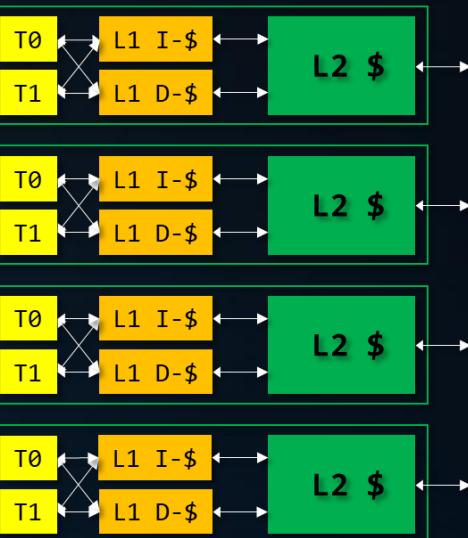
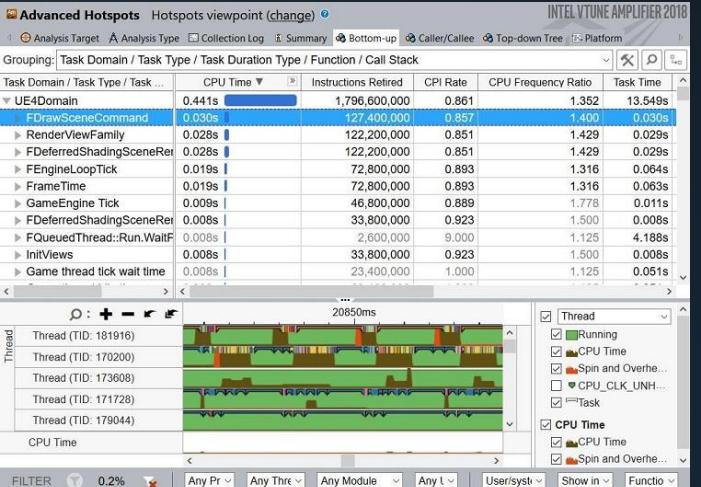
E * diffuse;
= true;

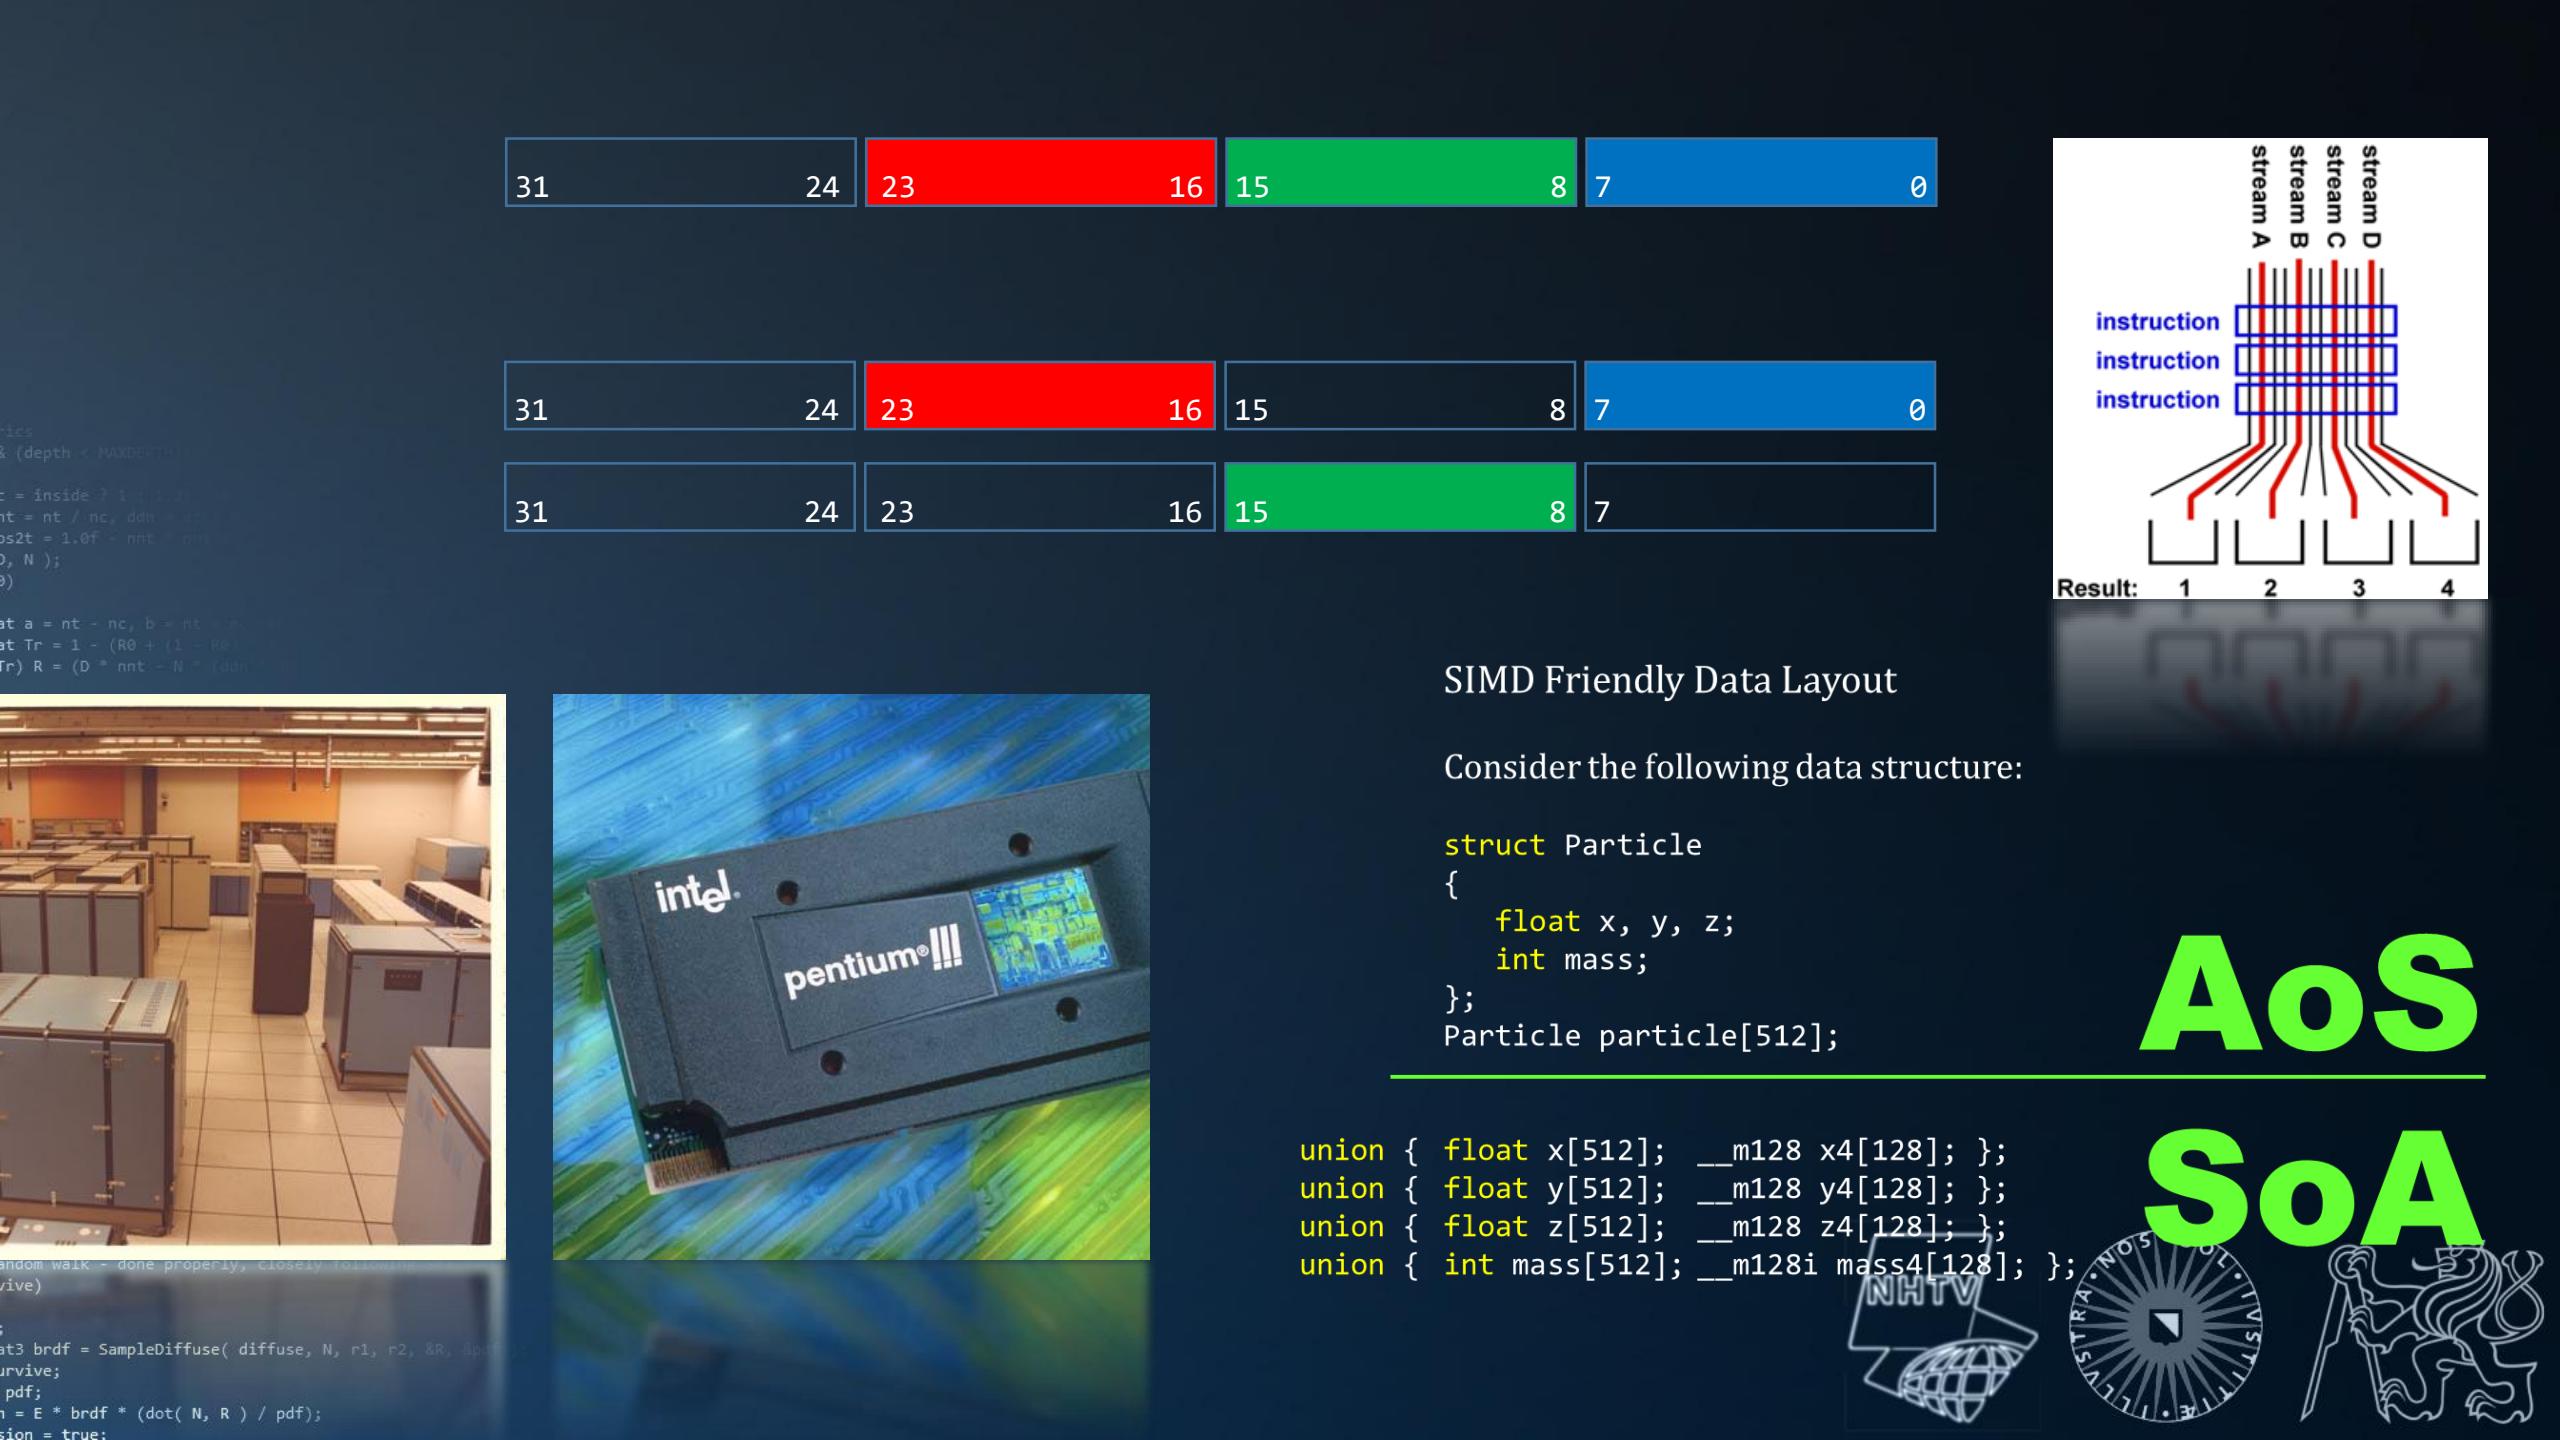
at x+=0, y = 0.1f;
signed int i = 0, j = 0x28929227;
if( kdiffuse &= 0, k < ITERATIONS; k++ )

// ...
x += y, y *= 1.01f;
survive = SurvivalProbability( diffuse );
estimation<= j, j ^= 0x17737352, i >>= 1, j /= 28763;
radiance = SampleLight( &rand, I, &L, &lighting );
x + radiance.y + radiance.z > 0) && (dot( N, r ) > 0)
if (with) x *= y;
true;
at brdfPdf = EvaluateDiffuse( I, &L, &lighting );
at3 factor = diffuse / INVPi;
at weight = Mis2( directPdf, brdfPdf );
at cosThetaOut = dot( N, r );
E * (weight * cosThetaOut / directPdf) * (radiance
random walk - done properly, closely following Smiley
alive);

at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf );
survive;
pdf;
n = E * brdf * (dot( N, R ) / pdf);
isalive = true;

```





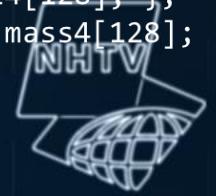
SIMD Friendly Data Layout

Consider the following data structure:

```
struct Particle
{
    float x, y, z;
    int mass;
};
```

Particle particle[512];

```
union { float x[512]; __m128 x4[128]; };
union { float y[512]; __m128 y4[128]; };
union { float z[512]; __m128 z4[128]; };
union { int mass[512]; __m128i mass4[128]; };
```



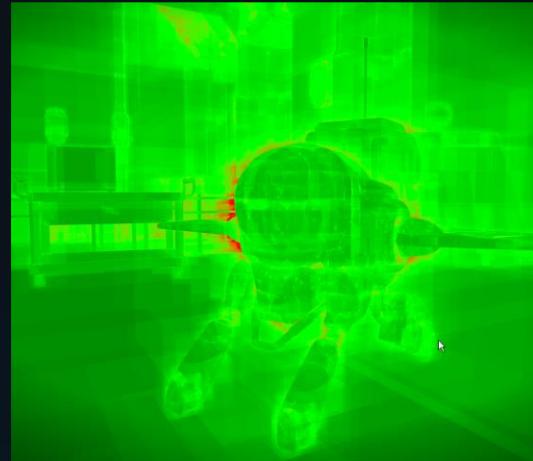
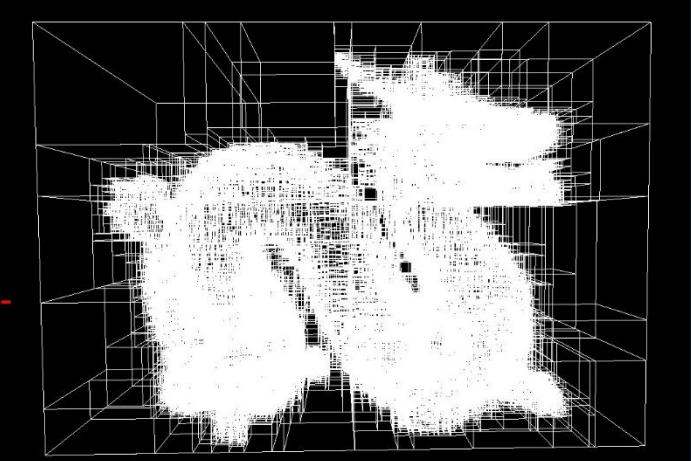
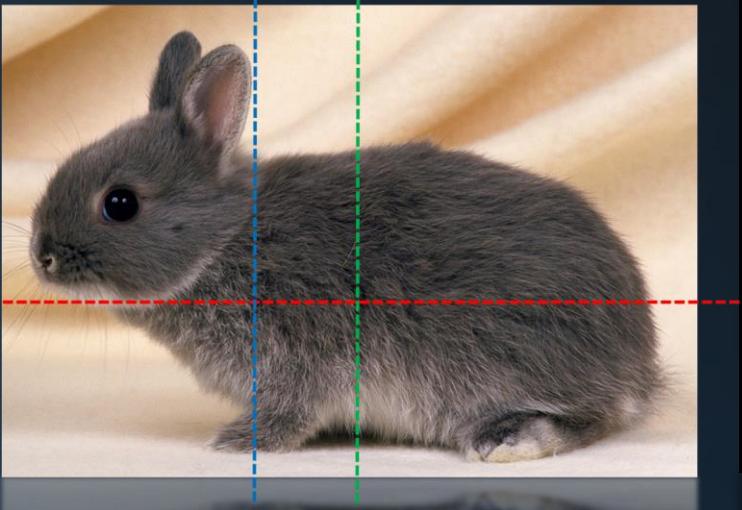
AoS

SoA

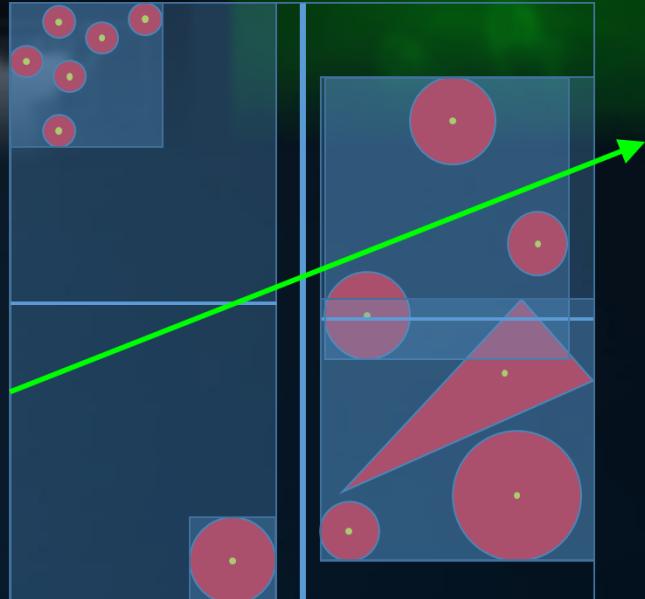
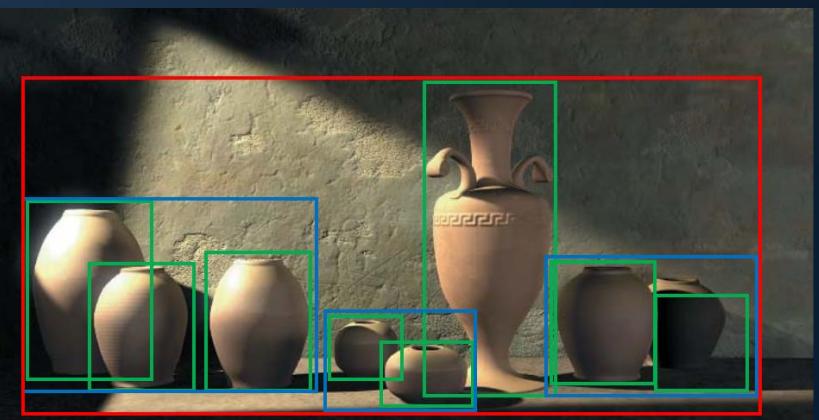
```
rics  
    & (depth < MAXDEPTH)  
  
    = inside ? 1 : 1.0f  
    nt = nt / nc; ddn = ddn / nc;  
    pos2t = 1.0f - nt * (1.0f - ddn);  
    D, N);  
}  
  
E * diffuse;  
= true;  
  
refl + refr)) && (depth < MAXDEPTH)  
    refl * E * diffuse;  
    = true;  
  
    D, N);  
    refl * E * diffuse;  
    = true;  
  
MAXDEPTH)
```

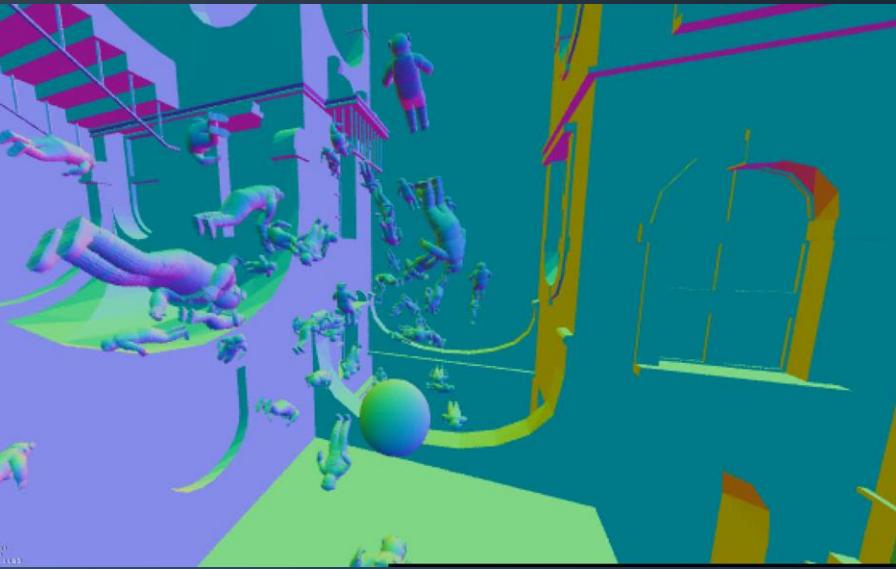
struct BVHNode

```
{  
    AABB bounds; // 24 bytes  
    int leftFirst; // 4 bytes  
    int count; // 4 bytes, total 32 ☺
```



```
survive = SurvivalProbability( diffuse  
estimation - doing it properly, closely  
if;  
radiance = SampleLight( &rand, I, &  
e.x + radiance.y + radiance.z ) < 0 ) && (dot( N, L ) >  
v = true;  
at brdfPdf = EvaluateDiffuse( L, N ) * Psurvive;  
at t3 factor = diffuse * INVPI;  
at weight = Mis2( directRdf, brdfPdf );  
at cosThetaOut = dot( N, L );  
E * (weight * cosThetaOut) / directRdf ) * (random  
walk - done properly, closely following S  
survive);  
at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf );  
survive;  
pdf;  
E * brdf * (dot( N, R ) / pdf);  
survive = true;
```





```
refl + refr) && (depth < MAXDEPTH)

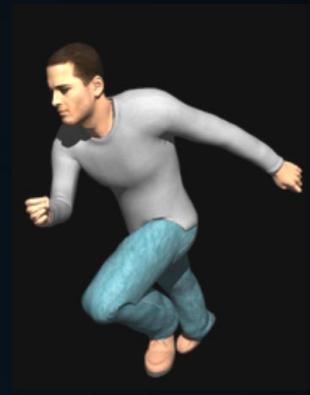
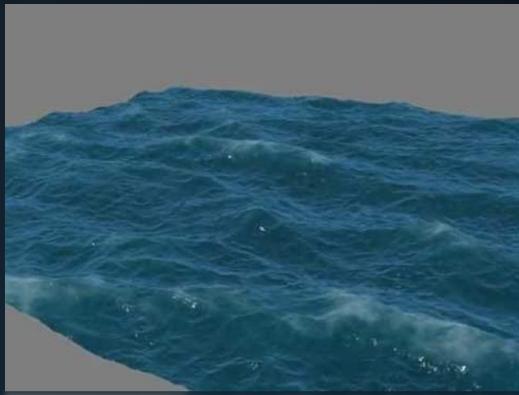
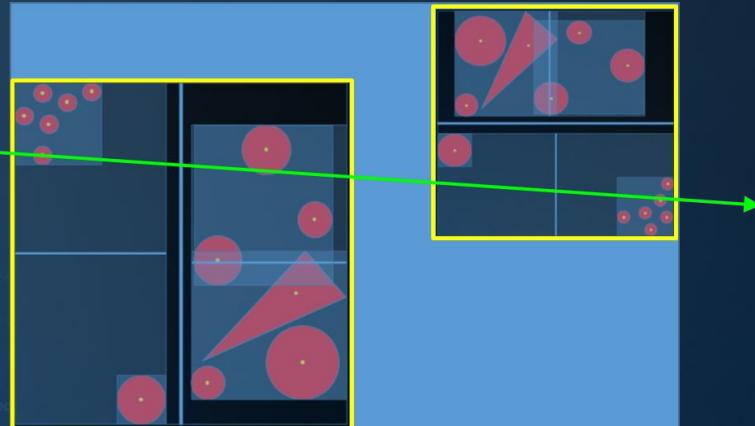
    N );
    refl * E * diffuse;
    = true;

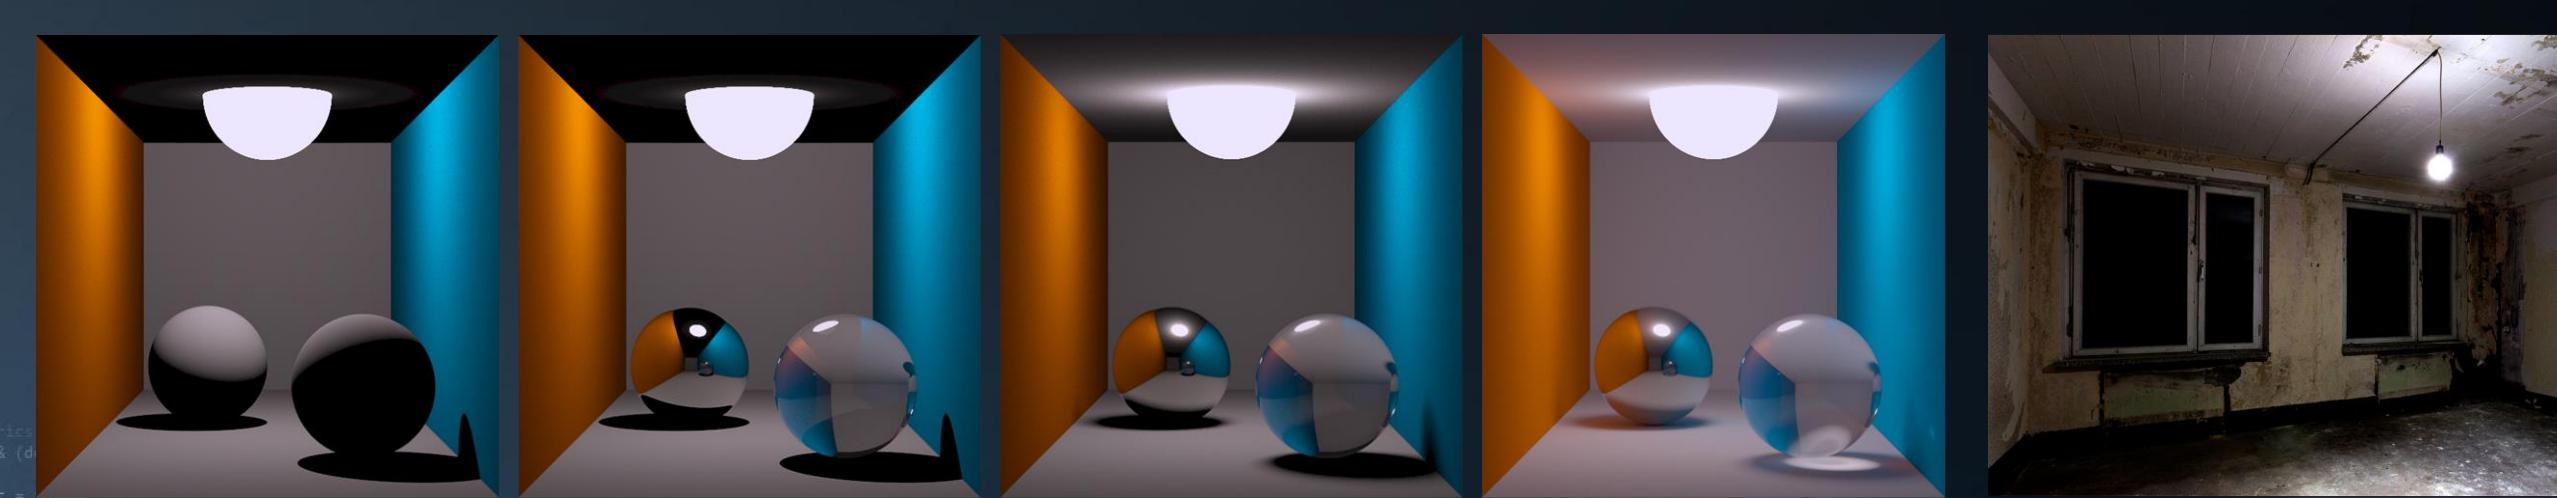
MAXDEPTH)

survive = SurvivalProbability( diffuse
estimation - doing it properly, closely
df;
radiance = SampleLight( &rand, I, &L, &lightbox,
x + radiance.y + radiance.z ) > 0 ) && depth
N );

    = true;
at brdfPdf = EvaluateDiffuse( L, N ) * psurvive;
at at3 factor = diffuse * INVPI;
at weight = Mis2( directPdf, brdfPdf );
at cosThetaOut = dot( N, L );
E * (weight * cosThetaOut) / directPdf ) * (radiance
random walk - done properly, closely following Smith
survive);

at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &p
survive;
pdf;
E * brdf * (dot( N, R ) / pdf);
ision = true;
```

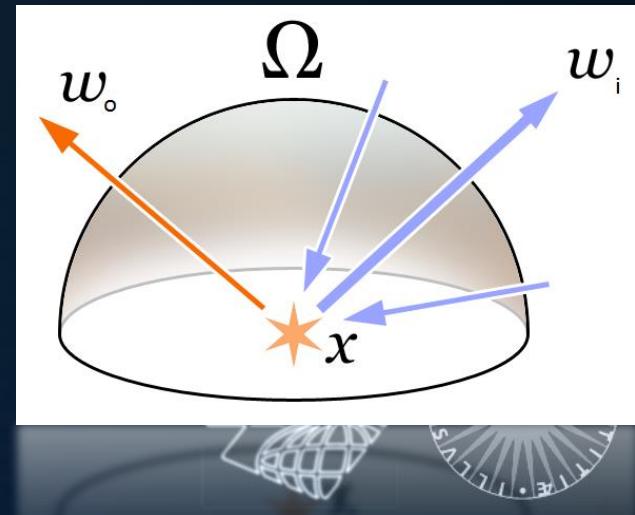
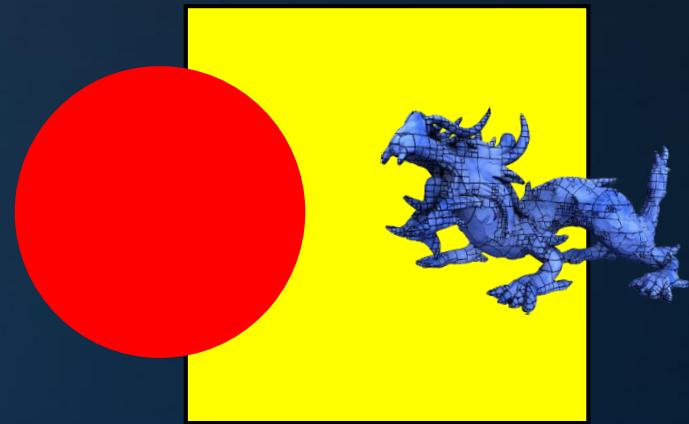
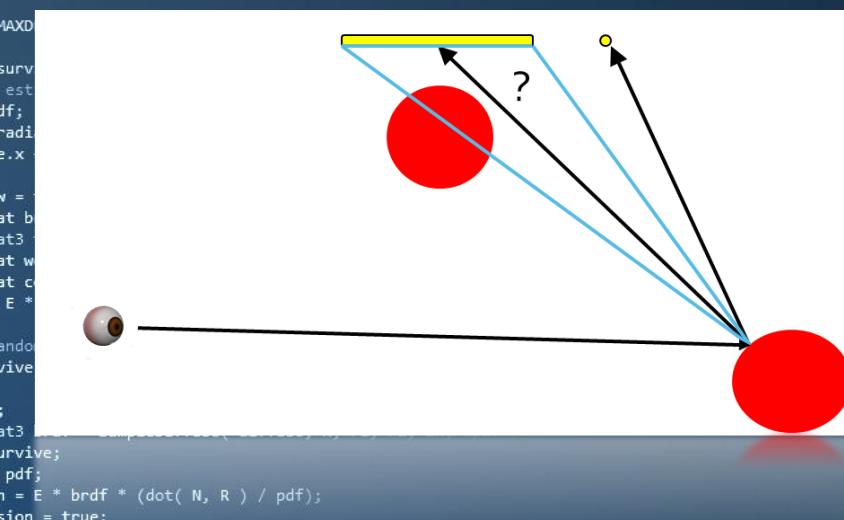


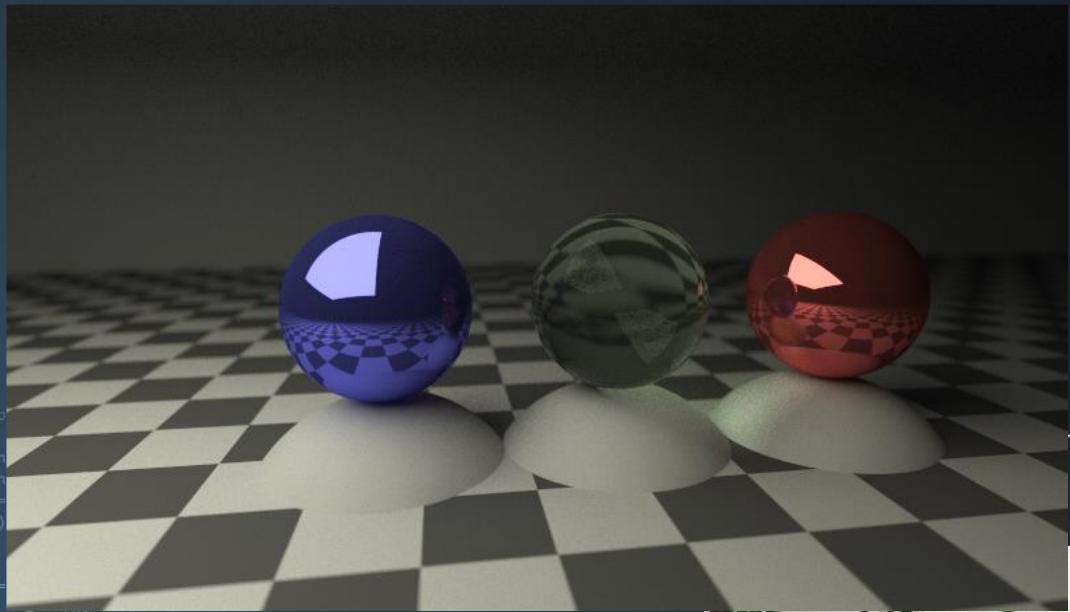


```
rics  
3 (d  
  
nt = nt / nc, dd  
pos2t = 1.0f < nnt ? pos2t : nnt / (nnt - 1.0f);  
(d  
N );  
  
at a = nt - nc, b = nt - nc  
at Tr = 1 - (R0 + (1 - R0) *  
Tr) R = (D * nnt - N * (dd  
E * diffuse;  
= true;  
refl + refr)) && (depth A_MAXIMUM_DEPTH  
0, N );  
refl * E * diffuse;  
= true;
```

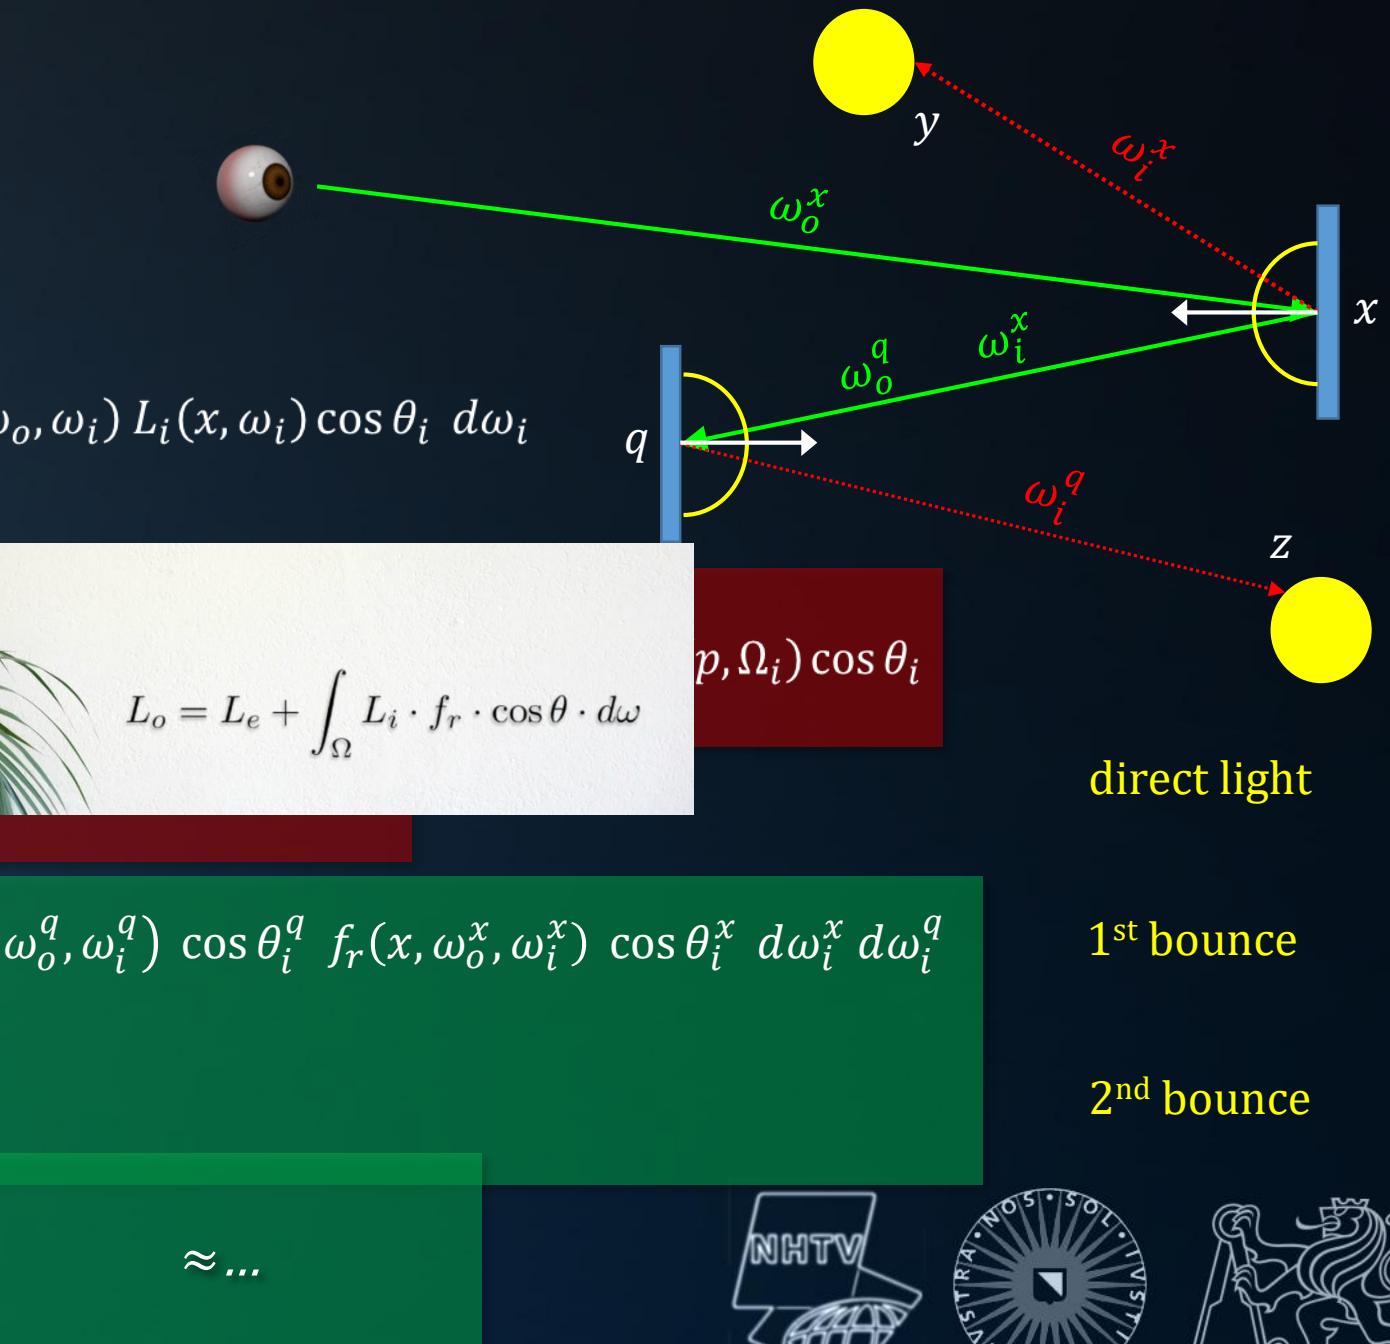
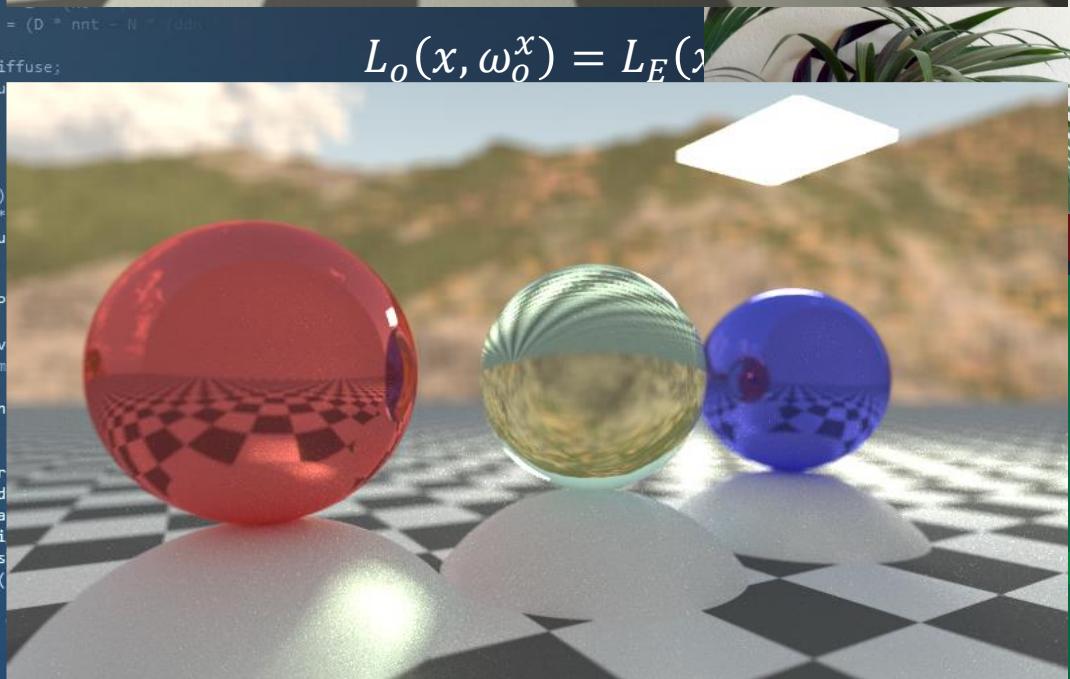
$$V_A = \int_A^B f(x) dx \approx \frac{B-A}{N} \sum_{i=1}^N f(X_i)$$

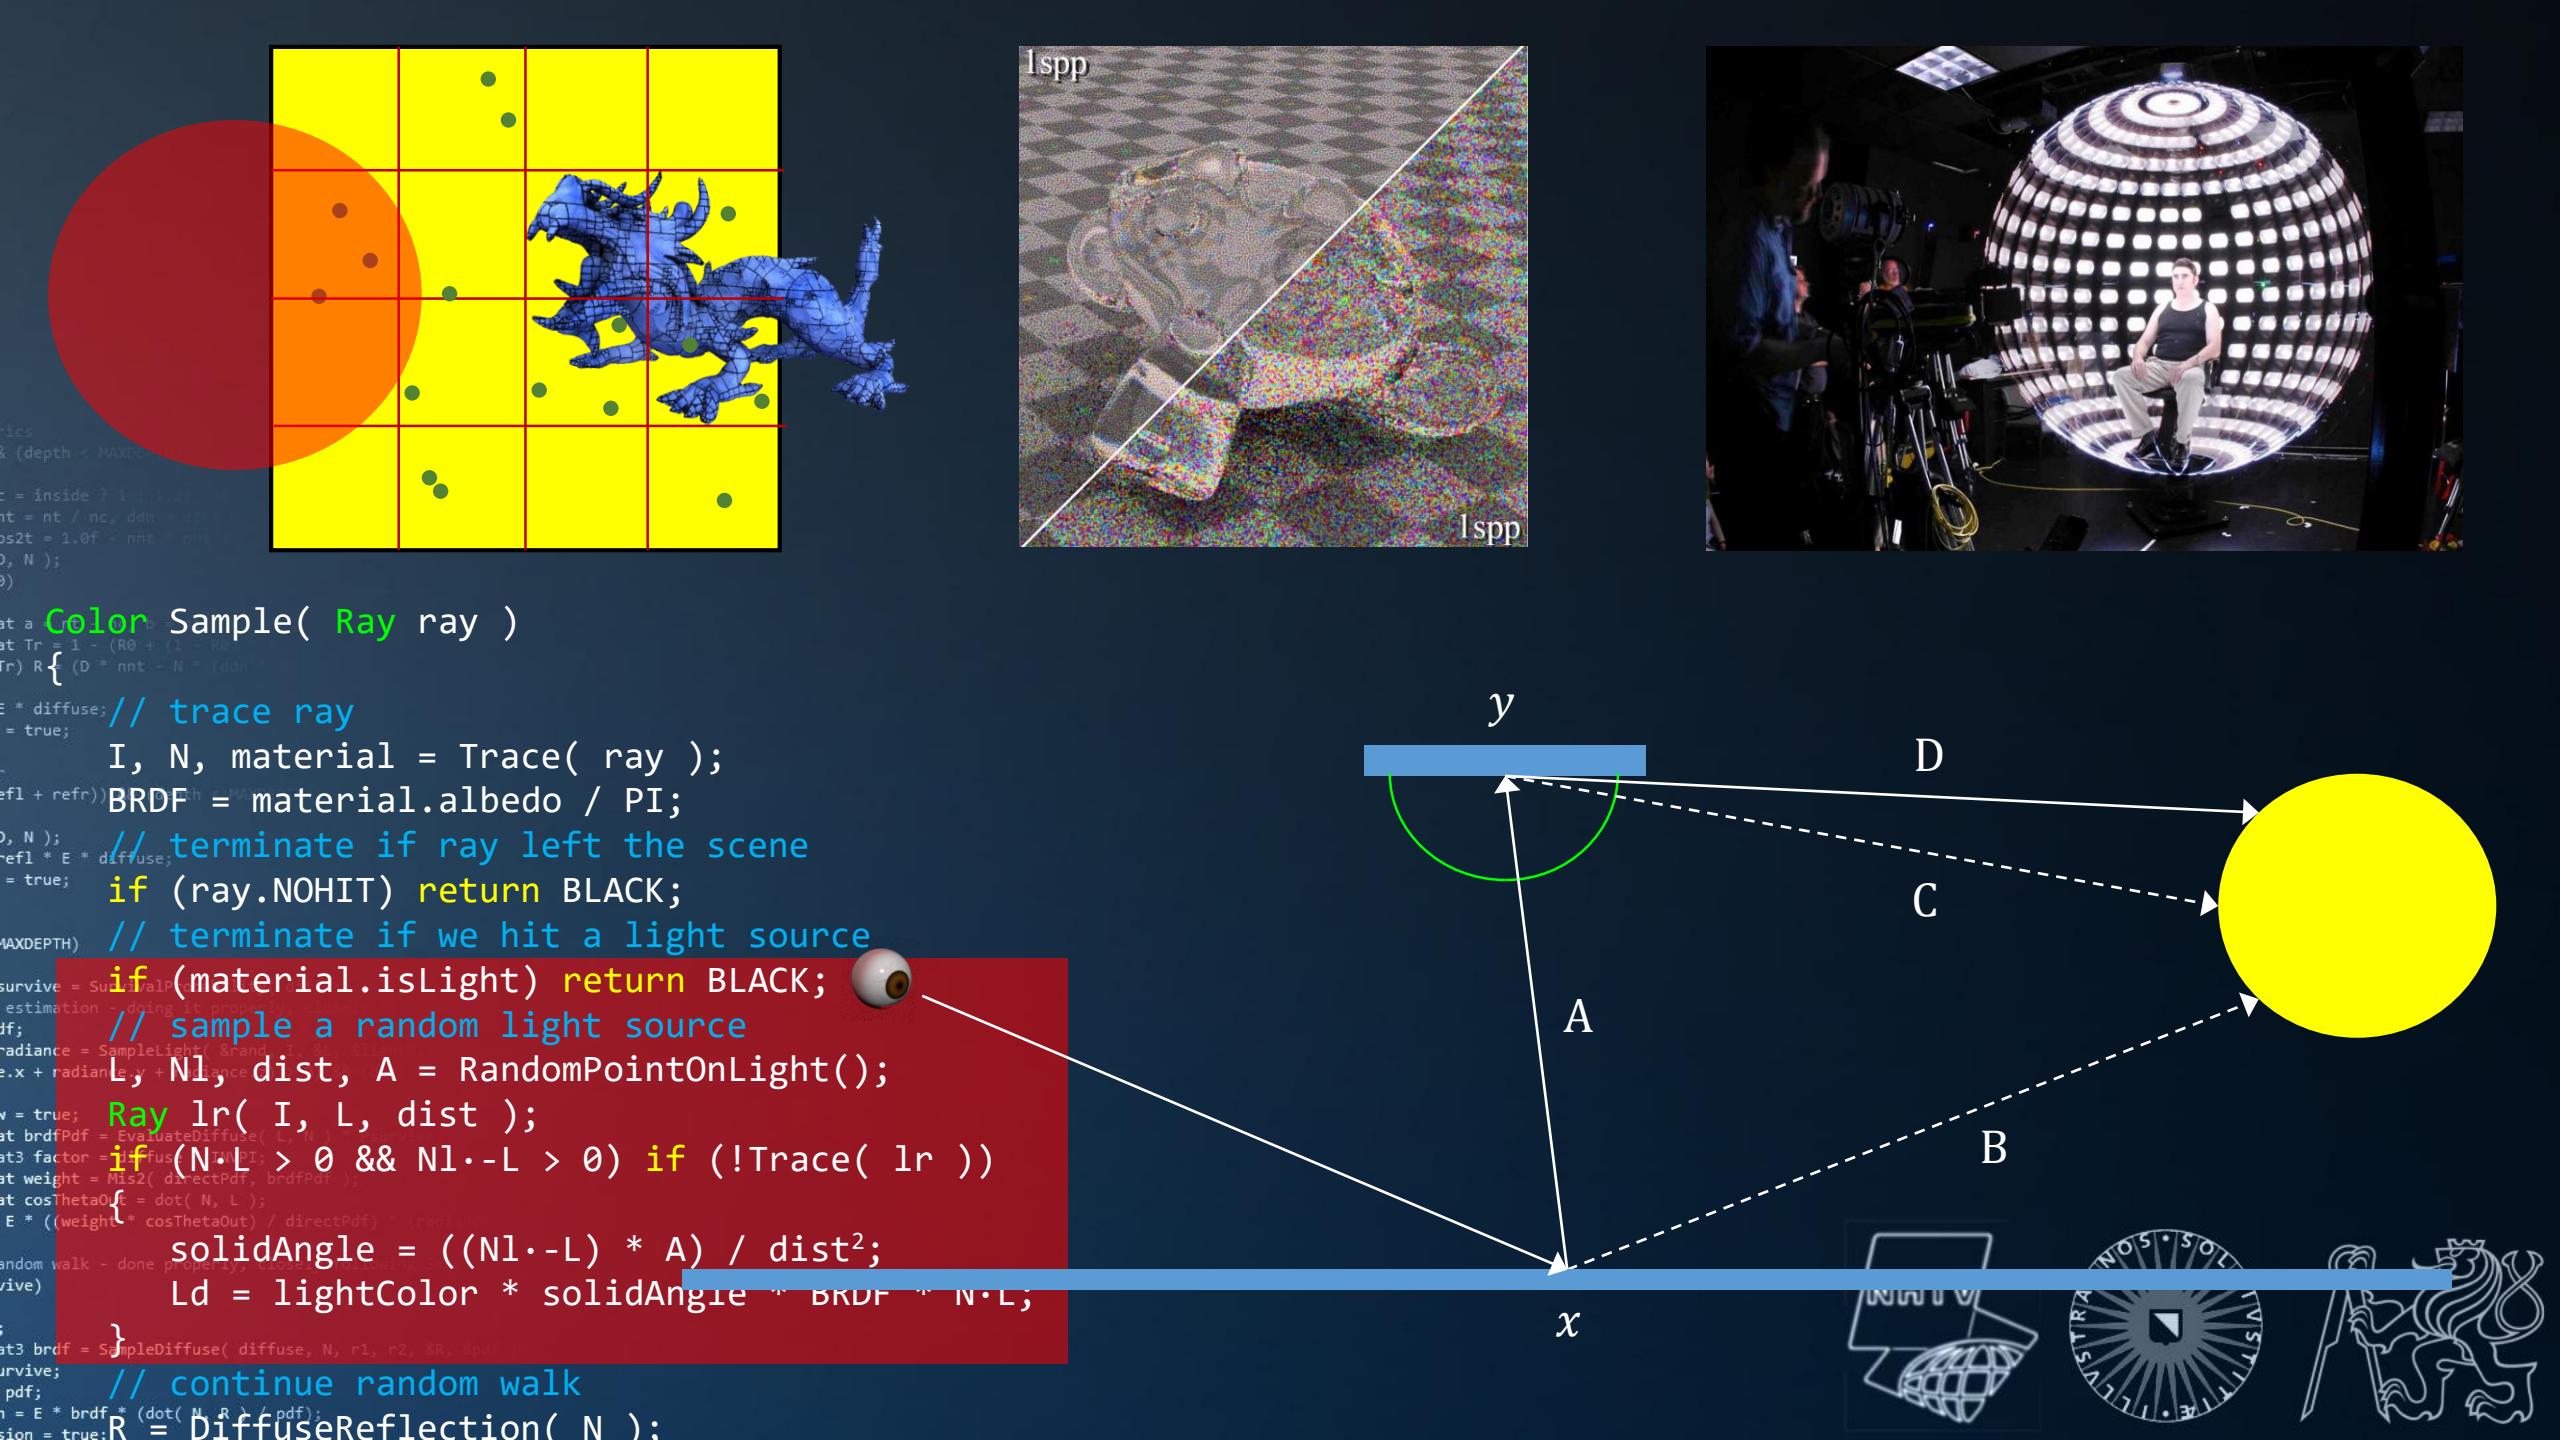
$$L(s \leftarrow x) = L_E(s \leftarrow x) + \int_A f_r(s \leftarrow x \leftarrow y) L(x \leftarrow y) G(x \leftrightarrow y) dA(y)$$





$$L_o(x, \omega_o^x) = L_E(x,$$

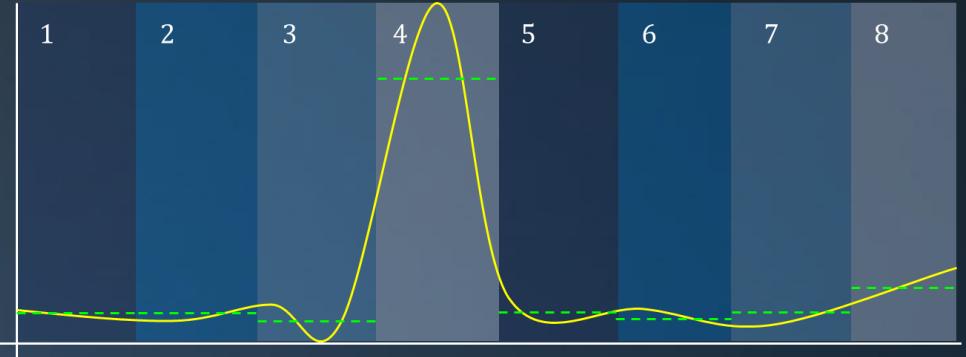




```

    Color Sample( Ray ray )
{
    I, N, material = Trace( ray );
    BRDF = material.albedo / PI;
    if (refl + refr) < 1.0f
        E * diffuse; // trace ray
    = true;
    if (ray.NOHIT) return BLACK;
    if (MAXDEPTH) // terminate if we hit a light source
        if (material.isLight) return BLACK;
        if; // sample a random light source
        radiance = SampleLight( &rand, &dist, &L, &Nl, &dist, &A = RandomPointOnLight());
        if (N.L > 0 && Nl.L > 0) if (!Trace( lr ))
            if (brdfPdf = EvaluateDiffuse( L, N ) > survive)
                if (at3 factor = Miss2( directPdf, brdfPdf );
                    weight = Miss2( directPdf, brdfPdf );
                    dot3 cosThetaOut = dot( N, L );
                    E * (weight * cosThetaOut) / directPdf ) > survive)
                        solidAngle = ((Nl.L) * A) / dist2;
                        Ld = lightColor * solidAngle * BRDF * N.L;
        if (survive)
            if (pdf);
            if (E * brdf * (dot( N, R ) / pdf));
            if (survive)
                R = DiffuseReflection( N );
}

```

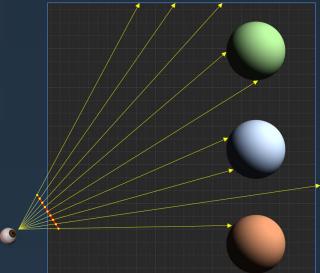
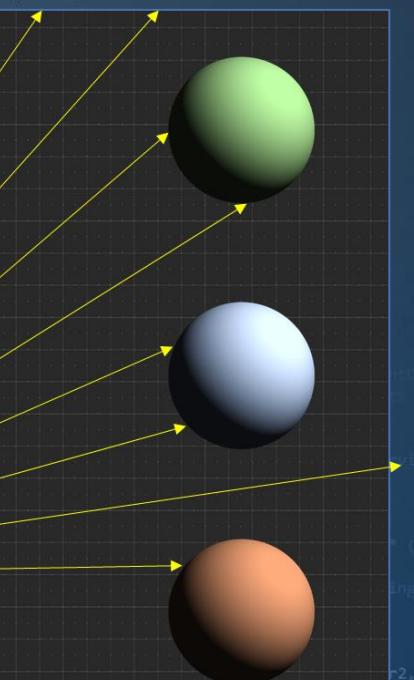
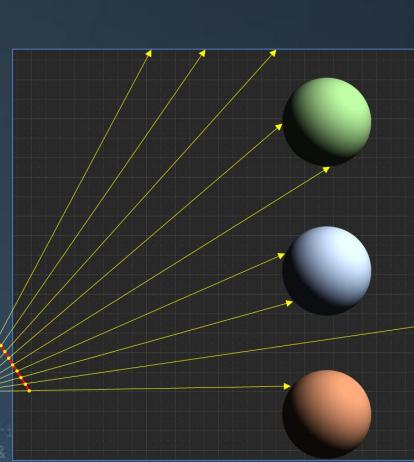


```

    Color Sample( Ray ray )
    {
        t a = nt - nc, b = nt + nc;
        t Tr = 1 - (R0 + (1-R0) * exp(-Nt));
        t R = (D * nnt - Nt * Nt) / (Nt * Nt);
        t E * diffuse;
        t = true;
        t refl + refr)) && (depth < MAXDEPTH)
        t refl * E * diffuse;
        t = true;
        t MAXDEPTH)
        R = CosineWeightedDiffuseReflection( N );
        survive = SurvivalProbability( d );
        estimation = doMonteCarlo( d );
        if( estimation > survive )
        {
            radiance = SampleLight( d );
            if( radiance.x + radiance.y + radiance.z > 0 ) &&
                radiance.x + radiance.y + radiance.z < 1 )
            {
                BRDF = material.albedo / PI;
                PDF = (N.R) / PI;
                Fi = Sample( r ) * (N.R) / PDF;
                E * (weight * cosineWeightedDiffuseReflection( N, L, ray, N, R, Tr, D, nnt, nt, nc, a, b, E, R, refl, refr, depth, maxDepth, survival, estimation, radiance, BRDF, PDF, Fi, weight, brdfPdf, brdfPd
            }
        }
    }

```





SCHWARZENEGGER

Get ready for the ride
of your life.

TOTAL RECAP

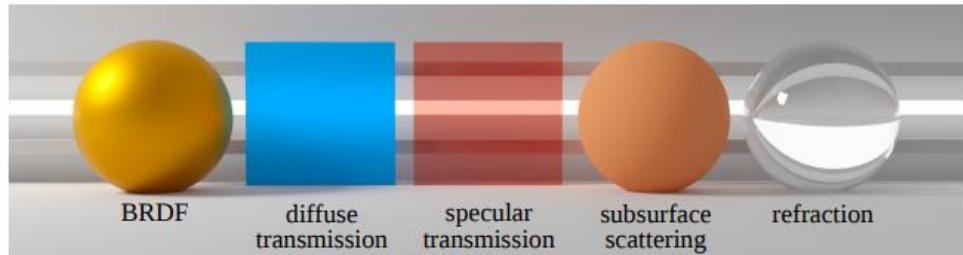
TOTAL RECAP

Agenda:

- Filtering
- Grand Recap
- Future Work



Materials



Extending the Disney BRDF to a BSDF with Integrated Subsurface Scattering

by Brent Burley, Walt Disney Animation Studios

1 Introduction

We introduced a new physically based shading model on Wreck-It Ralph [Bur12], and we used this single, general-purpose BRDF on all materials (except hair). This model, which has come to be known as the Disney BRDF, is able to reproduce a wide range of materials with only a few parameters. For our next film, Frozen, we continued to use this BRDF unmodified, but effects like refraction and subsurface scattering were computed separately from the BRDF, and indirect illumination was approximated using point clouds. All of these effects were combined through ad hoc shading in an additive way.

Starting with Big Hero 6 in 2014, we switched from ad hoc lighting and shading to path-traced global illumination where refraction, subsurface scattering, and indirect illumination were all integrated

Other resources:

Implementing the Disney BSDF
schuttejoe.github.io/post/disneybsdf

Efficient Rendering of Layered Materials using an Atomic Decomposition with Statistical Operators. Belcour, 2018.

Eric Heitz's Research Page:
eheitzresearch.wordpress.com/research
(all his work is good)



Spectral Rendering

Stratified Wavelength Clusters for Efficient Spectral Monte Carlo Rendering

Glenn F. Evans

Michael D. McCool

Computer Graphics Laboratory
Department of Computer Science
University of Waterloo

Abstract

Wavelength dependent Monte Carlo rendering can correctly and generally capture effects such as spectral caustics (rainbows) and chromatic aberration. It also improves the colour accuracy of reflectance models and of illumination effects such as colour bleeding and metamerism.

The stratified wavelength clustering (SWC) strategy carries several wavelength stratified radiance samples along each light transport path. The cluster is split into several paths or degraded into a single path only if a specular refraction at the surface of a dispersive material is encountered along the path. The overall efficiency of this strategy is high since the fraction of clusters that need to be split or degraded in a typical scene is low, and also because specular dispersion tends to decrease the source colour variance, offsetting the increased amortized cost of generating each path.

Key words: Monte Carlo methods, wavelength dependent (spectral) rendering, caustics, rainbows, refraction, reflectance, global illumination.

1 Introduction

The faults of using three component colour models for computing reflectance, absorption and dispersion are well known

per we demonstrate not only that this is feasible, but also that with a strategy we call *stratified wavelength clustering* (SWC) the marginal cost is negligible for Monte Carlo ray tracing and bidirectional path tracing.

2 Outline

Prior work on wavelength dependent and spectral rendering is surveyed in Section 3. Section 4 presents the stratified wavelength cluster strategy. Three splitting strategies are also presented and compared. In Section 5 a bidirectional path tracer that uses stratified wavelength clusters is described. Finally, in Section 6 we present our results, comparing and contrasting crude Monte Carlo integration over wavelength, quasi Monte Carlo integration over wavelength using Halton sequences, and the stratified wavelength cluster strategy.

3 Background

Before presenting our approach for extending Monte Carlo global illumination algorithms to wavelength-dependent (spectral) rendering, we first review the relationship of spectral power densities to perceived colour and review wavelength-dependent phenomena that have a significant effect on the generated image. Previous algorithms and representations for spectral ren-

Other resources:

Hero Wavelength Spectral Sampling. Wilie et al., 2014.

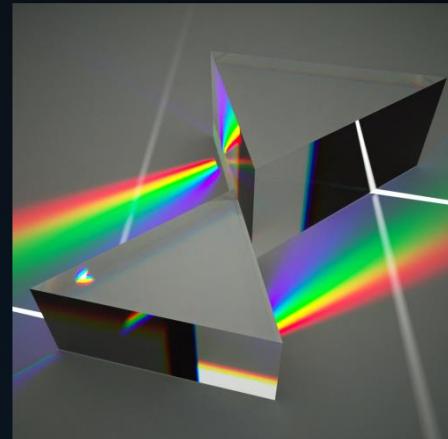
Physically Meaningful Rendering using Tristimulus Colours. Meng et al., 2015.

PBRT. Pharr et al., 2004-2018.

<https://www.pbrt.org>

Mitsuba Renderer. W. Jakob, 2014.

<https://www.mitsuba-renderer.org>



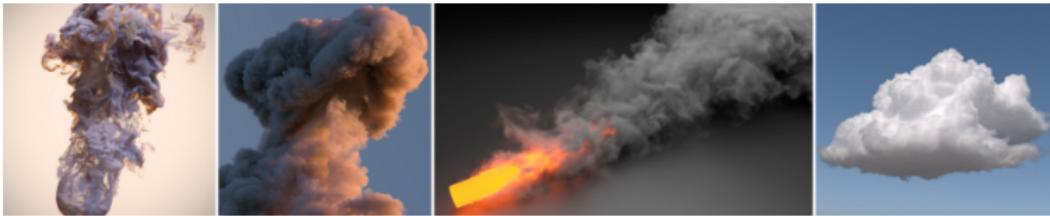
Participating Media

Monte Carlo methods for physically based volume rendering

Jan Novák^{1,2*} Iliyan Georgiev^{2,3*} Johannes Hanika^{3,2} Jaroslav Křivánek^{4,2} Wojciech Jarosz⁵

¹Disney Research ²Solid Angle ³Karlsruhe Institute of Technology ⁴Charles University, Prague ⁵Dartmouth College

In ACM SIGGRAPH Courses, 2018



Various media rendered with Monte Carlo methods for physically based simulation of light transport in volumes. The three images on the left are courtesy of Lee Griggs.

[Abstract](#) [Downloads](#) [Cite](#) [Related](#)

Abstract

We survey methods that utilize Monte Carlo (MC) integration to simulate light transport in scenes with participating media. The goal of this course is to complement a recent Eurographics 2018 state-of-the-art report providing a broad overview of most techniques developed to date, including a few methods from neutron transport, with a focus on concepts that are most relevant to CG practitioners.

The wide adoption of path-tracing algorithms in high-end realistic rendering has stimulated many diverse research initiatives aimed at efficiently rendering scenes with participating media. More computational power has enabled holistic approaches that tie volumetric effects and surface scattering together and simplify authoring workflows. Methods that were previously assumed to be incompatible have been unified to allow renderers to benefit from each method's respective strengths. Generally, investigations have shifted away from specialized solutions, e.g. for single- or multiple-scattering approximations or analytical methods, towards the more versatile Monte Carlo algorithms that are currently enjoying a widespread success in many production settings.

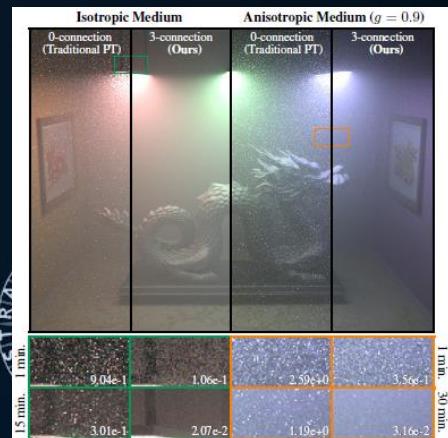
The goal of this course is to provide the audience with a deep, up-to-date understanding of key techniques for free-path sampling, transmittance estimation, and light-path construction in participating media, including those that are presently utilized in production rendering systems. We present a coherent overview of the fundamental building blocks and we contrast the various advanced methods that build on them, providing attendees with guidance for implementing existing solutions and developing new ones.

Downloads

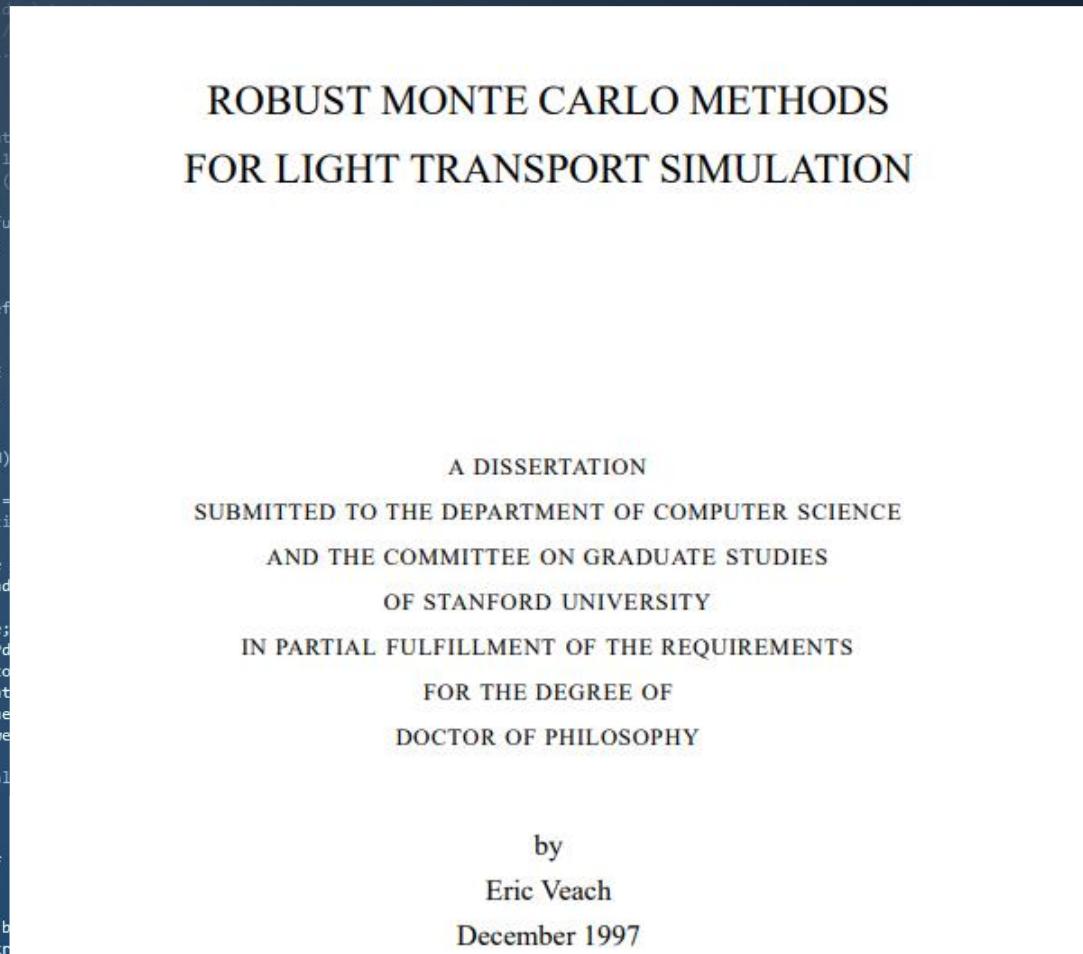
Other resources:

Importance Sampling Techniques for Path Tracing in Participating Media. Kulla and Fajardo, 2014.

Area Light Equi-Angular Sampling on ShaderToy:
<https://www.shadertoy.com/view/ldXGzS>



Light Transport



Other resources:

The Rendering Equation. James T. Kajiya, 1986.

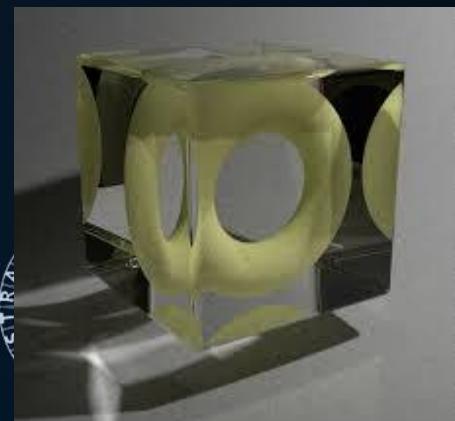
Bidirectional Path Tracing. Michal Vlnas, 2018 (student paper).

Global Illumination using Photon Maps. Jensen, 1996.

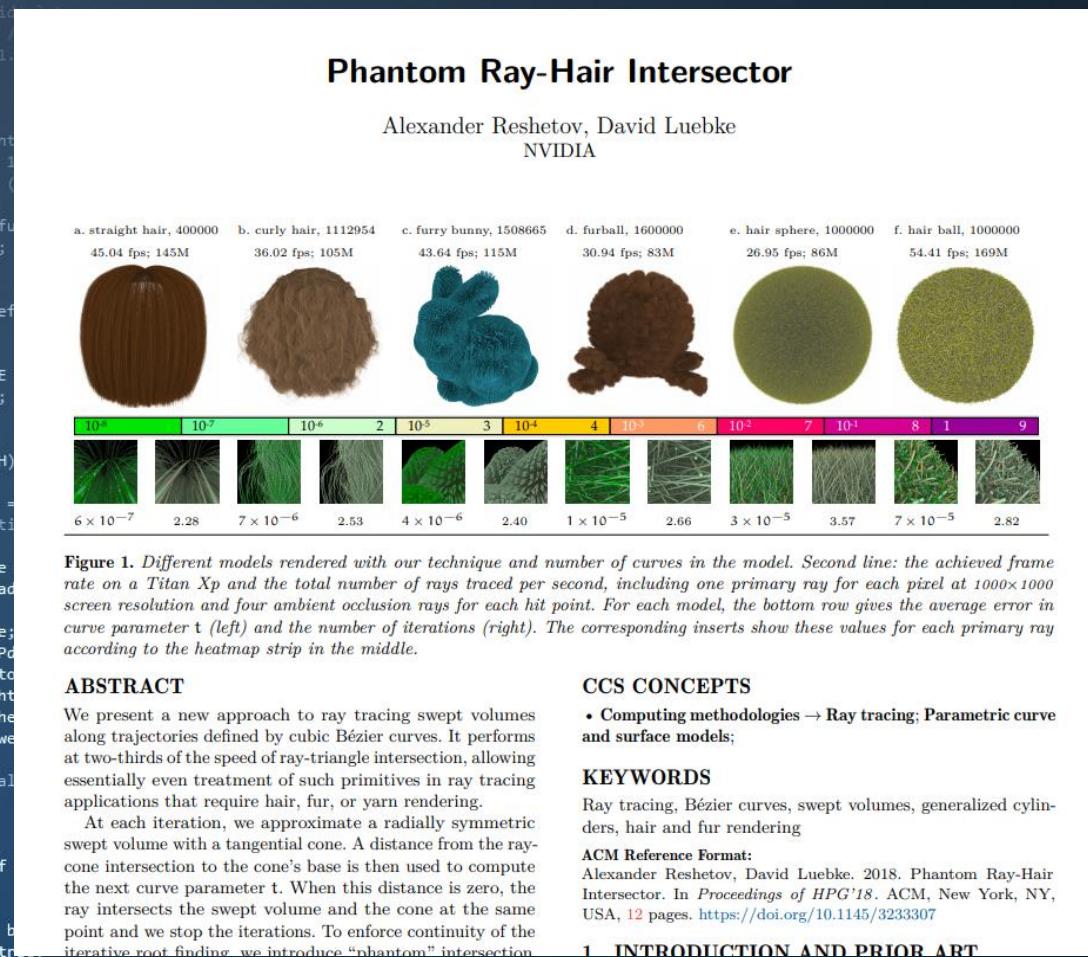
Progressive Photon Mapping. Hachisuka et al., 2008.

(article on) Vertex Connection and Merging, Georgev et al., 2012.

<https://schuttejoe.github.io/post/vertexconnectionandmerging>

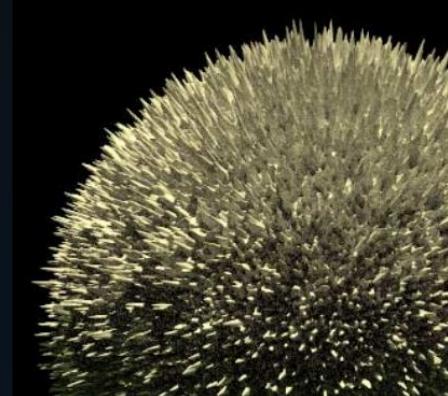


Primitives



Other resources:

Direct Ray Tracing of Smoothed and Displacement Mapped Triangles. Smits et al., 2000.



Direct Ray Tracing of Phong Tessellation. Ogaki and Tokuyoshi, 2011.



Two-Level Ray Tracing with Reordering for Highly Complex Scenes. Hanika et al., 2010.



Render Comparison



REDSHIFT 2.6.31



ARNOLD 5.2.2.0

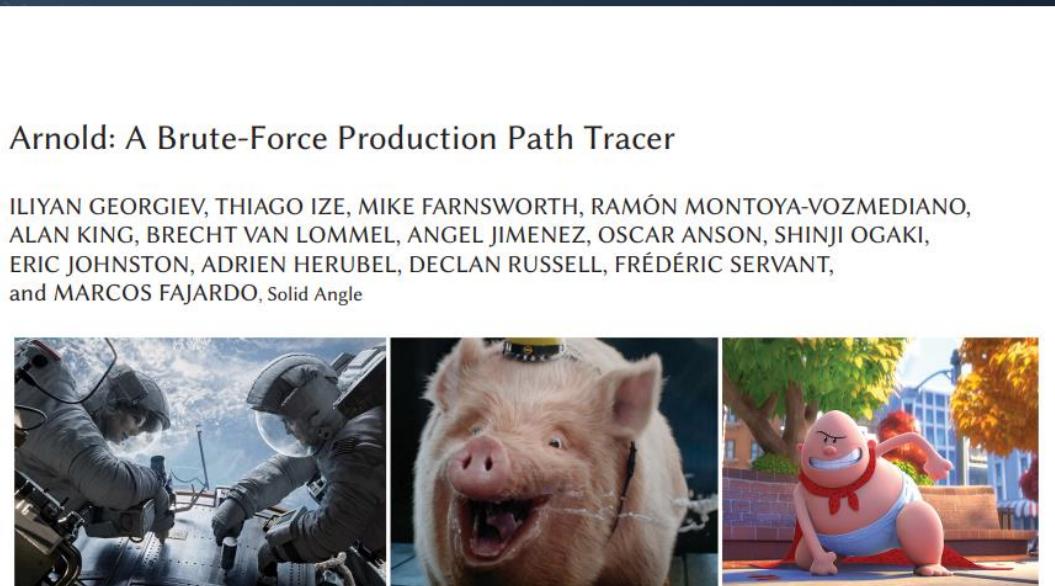


V-RAY NEXT, Beta 2, Build 4.11.01



HOUDINI 17.0.352 (MANTRA)

Production



Arnold: A Brute-Force Production Path Tracer

ILIAN GEORGIEV, THIAGO IZE, MIKE FARNSWORTH, RAMÓN MONToya-VOZMEDIANO, ALAN KING, BRECHT VAN LOMMEL, ANGEL JIMENEZ, OSCAR ANSON, SHINJI OGAKI, ERIC JOHNSTON, ADRIEN HERUBEL, DECLAN RUSSELL, FRÉDÉRIC SERVANT, and MARCOS FAJARDO, Solid Angle



Fig. 1. Arnold is a path-tracing renderer used for the production of photo-realistic and art-directed visual effects in feature films (left), commercials (middle), animated films (right), television series, music videos, game cinematics, motion graphics, and others. *Gravity* ©2013 Warner Bros. Pictures, courtesy of Framestore; *Racing Faces* ©2016 Opel Motorsport, courtesy of The Mill; *Captain Underpants* ©2017 DreamWorks Animation.

Arnold is a physically based renderer for feature-length animation and visual effects. Conceived in an era of complex multi-pass rasterization-based workflows struggling to keep up with growing demands for complexity and realism, Arnold was created to take on the challenge of making the simple and elegant approach of brute-force Monte Carlo path tracing practical for production rendering. Achieving this required building a robust piece of ray-tracing software that can ingest large amounts of geometry with detailed shading and lighting and produce images with high fidelity, while scaling well with the available memory and processing power.

Arnold's guiding principles are to expose as few controls as possible, provide rapid feedback to artists, and adapt to various production workflows. In this article, we describe its architecture with a focus on the design and implementation choices made during its evolution and development.

ACM Reference format:

Iliyan Georgiev, Thiago Ize, Mike Farnsworth, Ramón Montoya-Vozmediano, Alan King, Brecht Van Lommel, Angel Jimenez, Oscar Anson, Shinji Ogaki, Eric Johnston, Adrien Herubel, Declan Russell, Frédéric Servant, and Marcos Fajardo. 2018. Arnold: A Brute-Force Production Path Tracer. *ACM Trans. Graph.* 37, 3, Article 32 (July 2018), 12 pages.
<https://doi.org/10.1145/3182160>

1 INTRODUCTION

At a purely technical level, visual realism in computer-generated imagery boils down to two areas: (1) amount of scene detail, e.g., number of geometric primitives, amount of textures, variety of ma-

Other resources:

The Iray Light Transport Simulation and Rendering System. Keller et al., 2017.



The Design and Evolution of Disney's Hyperion Renderer. Burley et al., 2018.



Manuka: A batch-shading architecture for spectral path tracing in movie production. Fascione et al., 2018.



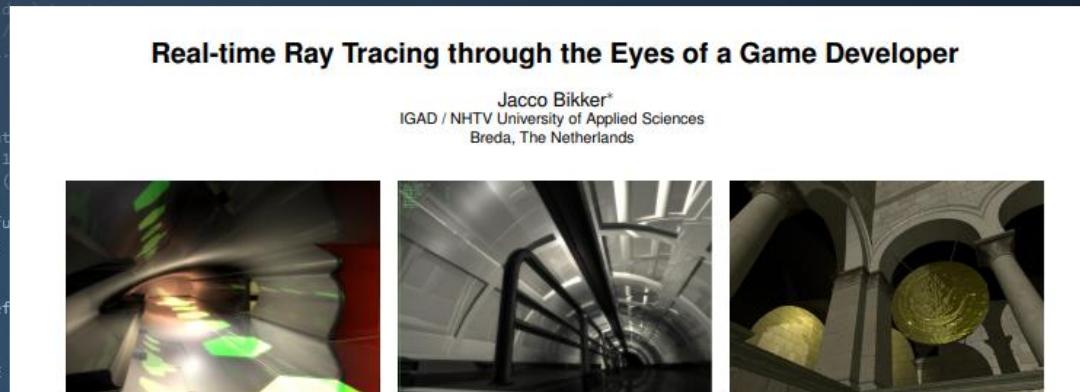
Mitsuba 2: A Retargetable Forward and Inverse Renderer. Nimier-David, 2019.



ALITA

BATTLE ANGEL

Games



Real-time Ray Tracing through the Eyes of a Game Developer
Jacco Bikker^{*}
IGAD / NHTV University of Applied Sciences
Breda, The Netherlands

Figure 1: Real-time ray traced images from our experimental ARAUNA engine, used in a student project.

ABSTRACT
There has been and is a tremendous amount of research on the topic of ray tracing, spurred by the relatively recent advent of real-time ray tracing and the inevitable appearance of consumer hardware capable of handling this rendering algorithm. Besides researchers, the prospect of a brave new world attracts hobbyists (such as demo coders) and game developers: Ray tracing promises an elegant and fascinating alternative to z-buffer approaches, as well as more intuitive graphics and games development. This article provides a view from the inside on ray tracing in games and demos, where the emphasis is on performance and short-term practical usability. It covers the way science is applied, the unique contribution of these developers to the field and their link with the research community.

The Arauna ray tracer, developed at the NHTV university of applied science, is used as an example of a ray tracer that has been specifically build with games and performance in mind. Its purpose and architecture, as well as some implementation details are presented.

Index Terms: I.3.7 [Computer Graphics]: Three-dimensional Graphics and Realism—Ray Tracing

1 INTRODUCTION

like Heaven7 [5] and the RealStorm benchmark [6] displayed the capabilities of this somewhat invisible demoscene movement [24]. The importance of ray tracing for games has also been recognised by researchers [17]. Ray tracing has been used extensively for offline rendering of game graphics (e.g. Myst and Riven [36]) and for lighting preprocessing as in e.g. Quake 2 [12]. Games are also frequently used as a test-case for real-time ray tracing: At the Breakpoint 2005 demo party, a fully playable real-time ray traced version of Quake 2 was shown by Wächter and Keller [14]. Ray traced versions of animated Quake 3 [19] and Quake 4 [20] walkthroughs where shown running on a PC cluster, and even the PS3 gameconsole is already being used for ray tracing [2].

Today, it has become clear that advancing ray tracing performance is not just a matter of better high-level algorithms: Low level optimization plays a crucial role. It is also clear that ray tracing performance is not dependent of fast ray traversal alone: Shading cost plays a significant role in real-time ray tracing [21, 31].

This article presents a description of the construction of the Arauna ray tracer, which was build as an exercise in applied science and with performance in mind. The Arauna ray tracer project started as an attempt to implement the ideas presented in Wald's PhD thesis on real-time ray tracing [30]. While previous attempts at interactive ray tracing used approximations (e.g., the Render Cache [34], the halodeck ray cache [15]), the OpenRT ray tracer [23]

Other resources:

OptiX: A General Purpose Ray Tracing Engine. Parker et al., 2010.

REAL-TIME RAYTRACING WITH NVIDIA RTX. Stich, 2018.

(not a lot of research so far...)

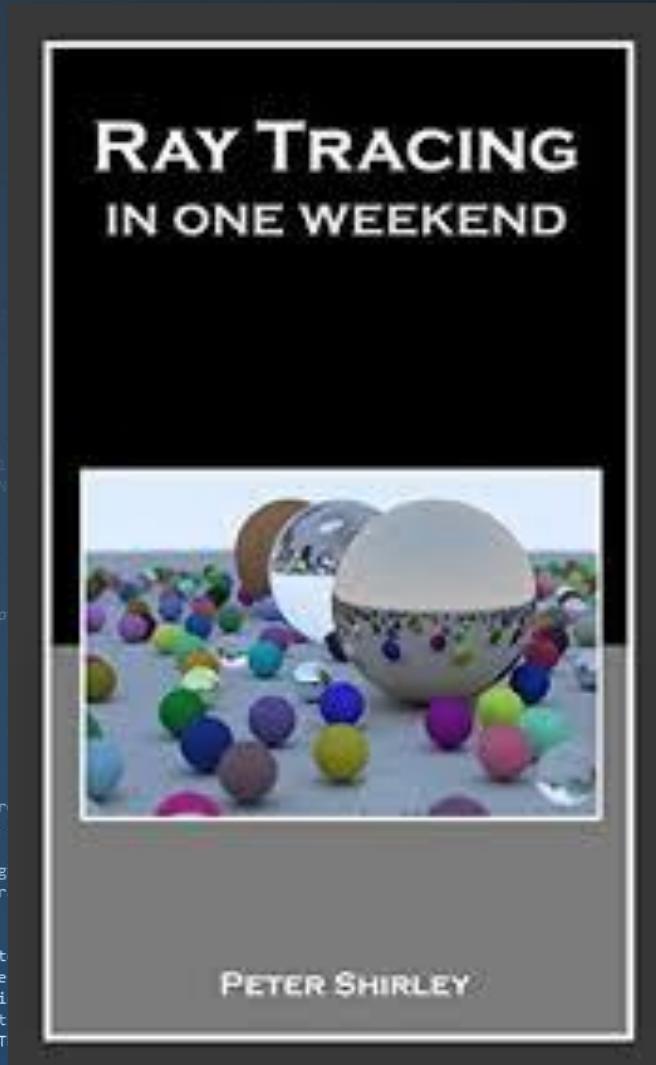


Massive Scenes

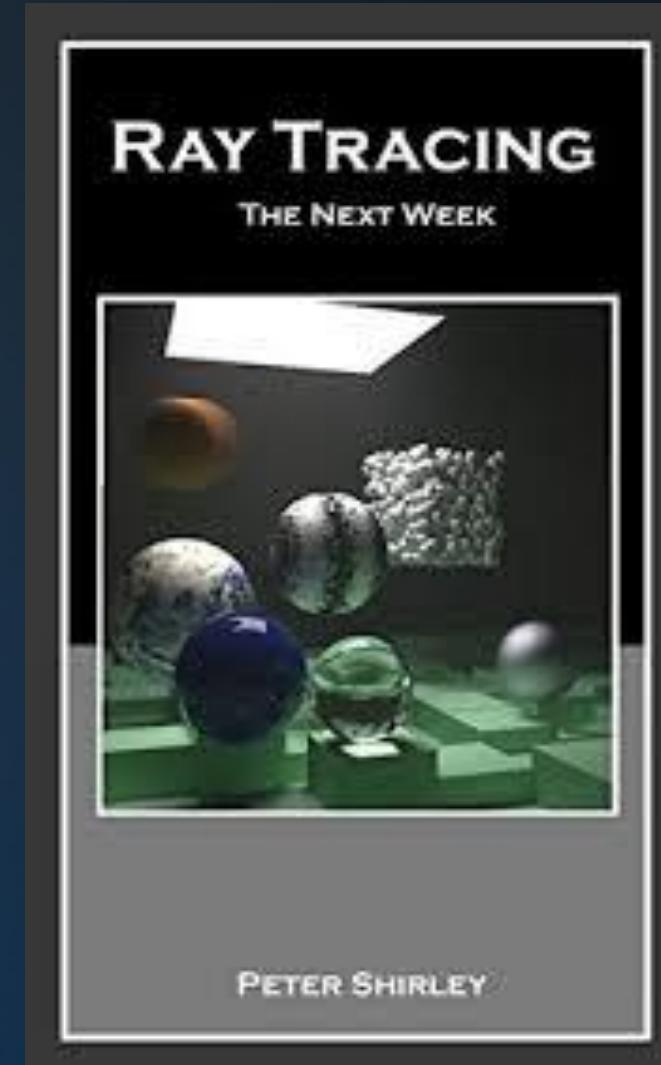


Moana Island scene. Heather Pritchett & Rasmus Tamstorf, 2016. ~ 15 billion primitives.





```
rics  
    & (depth < MAXDEPTH)  
    & (inside ? 1 : 1);  
    nt = nt / nc; ddn =  
    pos2t = 1.0f - nnt *  
    (D, N );  
}  
  
at a = nt - nc, b =  
at Tr = 1 - (R0 + (1  
- Tr) R = (D * nnt - N  
- E * diffuse;  
= true;  
  
-  
at refl + refr)) && (dep  
D, N );  
at refl * E * diffuse;  
= true;  
  
MAXDEPTH)  
  
survive = SurvivalPr  
estimation - doing  
if;  
radiance = SampleLig  
e.x + radiance.y + r  
  
v = true;  
at brdfPdf = Evaluat  
at3 factor = diffuse  
at weight = Mis2( di  
at cosThetaOut = dot  
E * ((weight * cost  
  
random walk - done properly, closely following S  
ive)  
  
;  
at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf )  
survive;  
pdf;  
n = E * brdf * (dot( N, R ) / pdf);  
ision = true;
```



Wednesday
13:00 – 17:00

course intro
LH2
template
Whitted
refactoring
RT-centric games

LAB 1

```
rics  
3 < depth < MAXDEPTH  
c = inside ? 1 : 1.2;  
nt = nt / ncy; ddn = ddn / ncy;  
os2t = 1.0f - nt * os2t;  
, N );  
)  
at a  
at Tr  
Tr ) R  
E * d  
= tr  
efl +  
, N  
refl  
= tr  
MAXDE  
surv  
estim  
df;  
radi  
e.x +  
v = t  
at br  
at3 f  
at we  
at co  
E *  
f) * (radian  
Psurvive  
light  
det  
)  
;  
at3 brdf = SampleDiffuse( diffuse, N, r1, r2, &R, &pdf  
survive;  
pdf;  
n = E * brdf * (dot( N, R ) / pdf);  
ision = true;
```

work @ home

End result day 2:
A solid Whitted-style
ray tracer, as a basis
for subsequent work.

LAB 1

Thursday
09:00 – 14:00

advanced Whitted
audio, AI & physics
faster Whitted
Heaven7

LAB 2



Friday
09:00 – 17:00

optimization
profiling, rules of
engagement
threading

LAB 3

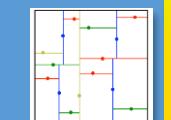


GLOBAL GAME JAM

Monday
09:00 – 17:00

acceleration
grid, BVH, kD-tree
SAH
binning

LAB 5



refitting
top-level BVH
threaded building

LAB 6

Tuesday
09:00 – 17:00

Monte-Carlo
Cook-style
glossy, AA
area lights, DOF

LAB 7



path tracing

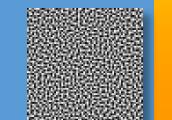
LAB 8



Thursday
09:00 – 17:00

random numbers
stratification
blue noise

LAB 9



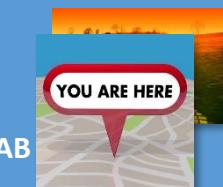
importance
sampling
next event
estimation

LAB 10

Friday
09:00 – 17:00

future work

LAB 10



path guiding

LAB 10



End result day 3:
A 5x faster tracer.

End result day 4:
A real-time tracer.

End result day 5:
Cook or Kajiya.

End result day 6:
Efficiency.

End result day 6:
Great product.



End of PART 11.

